

02_story_generation

June 14, 2025

1 Story Generation Model Analysis

1.1 Academic Analysis of Large Language Models for Children's Bedtime Story Generation

Author: University of London Student

Project: Smart Visual Storyteller for Children

Analysis Date: June 2025

1.1.1 Abstract

This notebook presents a comprehensive analysis of ten state-of-the-art large language models for generating children's bedtime stories. We evaluate OpenAI GPT-4o/GPT-4o-mini, Anthropic Claude-3.5-sonnet/Claude-3.5-haiku, Google Gemini-2.0-flash/Gemini-1.5-pro, and DeepSeek deepseek-chat/deepseek-v3 models across multiple dimensions including cost efficiency, execution speed, content quality, and age-appropriateness.

1.1.2 Research Objectives

1. **Performance Evaluation:** Assess computational efficiency and cost-effectiveness
2. **Quality Analysis:** Evaluate story quality, structure, and appropriateness
3. **Comparative Assessment:** Rank models for optimal story generation
4. **Implementation Guidance:** Provide evidence-based model selection recommendations

1.2 0. Requirements Installation

Before running this analysis, ensure all required packages are installed. Run the following cell if you haven't installed the requirements yet:

```
[66]: # Install required packages for analysis
      # Uncomment and run if packages are missing:

      # %pip install pandas matplotlib seaborn numpy scipy scikit-learn jupyter

      # Or install from requirements file:
      # %pip install -r ../requirements.txt

      # Check if all packages are available
      try:
```

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.preprocessing import StandardScaler
print(" All required packages are installed and available")
except ImportError as e:
    print(f" Missing package: {e}")
    print("Please uncomment and run the pip install commands above")

```

All required packages are installed and available

1.3 1. Environment Setup and Data Loading

```

[67]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path
import json
from scipy import stats
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
import warnings
warnings.filterwarnings('ignore')

# Configure plotting parameters for academic presentation
plt.style.use('default')
sns.set_palette("husl")
plt.rcParams['figure.figsize'] = (12, 8)
plt.rcParams['font.size'] = 12
plt.rcParams['axes.titlesize'] = 14
plt.rcParams['axes.labelsize'] = 12
plt.rcParams['xtick.labelsize'] = 10
plt.rcParams['ytick.labelsize'] = 10
plt.rcParams['legend.fontsize'] = 10

print("Analysis environment initialized successfully")
print(f"pandas version: {pd.__version__}")
print(f"numpy version: {np.__version__}")

```

Analysis environment initialized successfully
pandas version: 2.3.0
numpy version: 2.2.6

```

[68]: # Load the most recent story generation results
results_dir = Path('../..//results/story_generation')

```

```

csv_files = list(results_dir.glob('story_generation_results_*.csv'))

if not csv_files:
    raise FileNotFoundError("No story generation results found. Please run_
↳02_story_generation_collect.py first.")

# Use the most recent results file
latest_file = max(csv_files, key=lambda x: x.stat().st_mtime)
print(f"Loading data from: {latest_file.name}")

# Load the dataset
df = pd.read_csv(latest_file)

print(f"Dataset loaded successfully")
print(f"Total records: {len(df)}")
print(f"Unique models: {df['story_model'].nunique()}")
print(f"Unique images: {df['image_file'].nunique()}")
print(f>Date range: {latest_file.stat().st_mtime}")

```

```

Loading data from: story_generation_results_20250613_231958.csv
Dataset loaded successfully
Total records: 240
Unique models: 15
Unique images: 16
Date range: 1749847628.1840665

```

```

[69]: # Data exploration and basic information
print("=== DATASET CHARACTERIZATION ===\n")

print("Story generation models evaluated:")
for model in df['story_model'].unique():
    print(f" - {model}")

print(f"\nTest images:")
for image in df['image_file'].unique():
    print(f" - {image}")

print(f"\nData completeness assessment:")
expected_combinations = len(df['story_model'].unique()) * len(df['image_file'].
↳unique())
print(f" - Expected combinations: {len(df['story_model'].unique())} models ×_
↳{len(df['image_file'].unique())} images = {expected_combinations}")
print(f" - Actual records: {len(df)}")
print(f" - Missing records: {expected_combinations - len(df)}")

# Data quality assessment
print(f"\nData quality assessment:")

```

```

print(df.isnull().sum())

# Display dataset structure
print(f"\nDataset columns:")
print(list(df.columns))
print(f"\nSample data:")
df.head()

```

=== DATASET CHARACTERIZATION ===

Story generation models evaluated:

- gpt-4o
- gpt-4o-mini
- claude-opus-4
- claude-sonnet-4
- claude-3.7-sonnet
- claude-3.5-sonnet
- claude-3.5-haiku
- gemini-2.0-flash
- gemini-2.0-flash-lite
- gemini-1.5-pro
- gemini-1.5-flash
- mistral-large-latest
- mistral-medium-latest
- mistral-small-latest
- deepseek-chat

Test images:

- toy_01.jpeg
- toy_02.jpeg
- toy_03.jpeg
- toy_04.jpeg
- toy_05.jpeg
- toy_06.jpeg
- toy_07.jpeg
- toy_08.jpeg
- toy_09.jpeg
- toy_10.jpeg
- drawing_01.jpeg
- drawing_02.jpeg
- drawing_03.jpeg
- drawing_04.jpeg
- drawing_05.jpeg
- drawing_06.jpeg

Data completeness assessment:

- Expected combinations: 15 models × 16 images = 240
- Actual records: 240

- Missing records: 0

Data quality assessment:

image_file	0
image_type	0
image_caption	0
story_model	0
generated_story	0
execution_time	0
cost	0
word_count	0
quality_score	0
meets_length_req	0
has_title	0
contains_dialogue	0
positive_tone	0
story_structure	0
age_appropriate	0
bedtime_suitable	0

dtype: int64

Dataset columns:

['image_file', 'image_type', 'image_caption', 'story_model', 'generated_story', 'execution_time', 'cost', 'word_count', 'quality_score', 'meets_length_req', 'has_title', 'contains_dialogue', 'positive_tone', 'story_structure', 'age_appropriate', 'bedtime_suitable']

Sample data:

```
[69]:
```

	image_file	image_type	image_caption	\
0	toy_01.jpeg	toy	A doll with curly blonde hair and a red jumpsu...	
1	toy_01.jpeg	toy	A doll with curly blonde hair and a red jumpsu...	
2	toy_01.jpeg	toy	A doll with curly blonde hair and a red jumpsu...	
3	toy_01.jpeg	toy	A doll with curly blonde hair and a red jumpsu...	
4	toy_01.jpeg	toy	A doll with curly blonde hair and a red jumpsu...	

	story_model	generated_story	\
0	gpt-4o	Title: Lily's Cozy Adventure\n\nOnce upon a ti...	
1	gpt-4o-mini	Title: Daisy's Cozy Adventure\n\nOnce upon a t...	
2	claude-opus-4	Title: Rosie's Couch Adventure\n\nRosie the do...	
3	claude-sonnet-4	**Title: Rosie's Special Day**\n\nRosie the do...	
4	claude-3.7-sonnet	# The Couch Adventure\n\nMolly the doll with c...	

	execution_time	cost	word_count	quality_score	meets_length_req	\
0	4.52	0.004670	201	26.00	True	
1	4.70	0.000172	185	24.00	True	
2	11.63	0.018480	179	23.25	True	

3	6.86	0.003681	169	22.00	True
4	6.14	0.003981	191	24.62	True

	has_title	contains_dialogue	positive_tone	story_structure	\
0	True	True	True	True	
1	True	True	True	True	
2	True	True	True	True	
3	True	True	True	True	
4	False	True	True	True	

	age_appropriate	bedtime_suitable
0	True	True
1	True	True
2	True	True
3	True	True
4	True	True

1.4 2. Performance Metrics Analysis

```
[71]: # Performance statistics by model
print("=== PERFORMANCE METRICS ANALYSIS ===\n")

# Aggregate performance statistics by model
performance_stats = df.groupby('story_model').agg({
    'execution_time': ['mean', 'std', 'min', 'max'],
    'cost': ['mean', 'std', 'min', 'max']
}).round(4)

performance_stats.columns = ['_'.join(col) for col in performance_stats.columns]
print("Performance Statistics by Model:")
print(performance_stats)

# Calculate efficiency metrics
print("\n=== EFFICIENCY METRICS ===")
efficiency_metrics = df.groupby('story_model').agg({
    'execution_time': 'mean',
    'cost': 'mean'
}).round(4)

# Add cost per second metric
efficiency_metrics['cost_per_second'] = (efficiency_metrics['cost'] /
    efficiency_metrics['execution_time']).round(6)
efficiency_metrics['speed_rank'] = efficiency_metrics['execution_time'].rank()
efficiency_metrics['cost_rank'] = efficiency_metrics['cost'].rank()

print("Efficiency Rankings:")
print(efficiency_metrics.sort_values('execution_time'))
```

=== PERFORMANCE METRICS ANALYSIS ===

Performance Statistics by Model:

	execution_time_mean	execution_time_std \
story_model		
claude-3.5-haiku	7.3294	1.3386
claude-3.5-sonnet	6.5094	0.4776
claude-3.7-sonnet	7.1412	0.5581
claude-opus-4	11.7831	0.8059
claude-sonnet-4	7.5962	0.4994
deepseek-chat	13.1644	1.1027
gemini-1.5-flash	2.0706	0.1248
gemini-1.5-pro	5.9519	0.5986
gemini-2.0-flash	2.1431	0.3390
gemini-2.0-flash-lite	1.9762	0.1060
gpt-4o	6.6425	1.5734
gpt-4o-mini	5.2044	1.5751
mistral-large-latest	11.3275	5.3297
mistral-medium-latest	6.4062	1.8114
mistral-small-latest	3.5500	1.4485

	execution_time_min	execution_time_max	cost_mean \
story_model			
claude-3.5-haiku	5.42	10.26	0.0010
claude-3.5-sonnet	5.83	7.36	0.0040
claude-3.7-sonnet	6.14	7.88	0.0042
claude-opus-4	10.17	12.84	0.0190
claude-sonnet-4	6.86	8.47	0.0038
deepseek-chat	11.32	14.78	0.0003
gemini-1.5-flash	1.84	2.26	0.0003
gemini-1.5-pro	4.68	6.95	0.0043
gemini-2.0-flash	1.74	3.01	0.0003
gemini-2.0-flash-lite	1.84	2.19	0.0004
gpt-4o	4.27	9.33	0.0049
gpt-4o-mini	2.15	8.17	0.0002
mistral-large-latest	5.08	22.78	0.0020
mistral-medium-latest	4.32	11.42	0.0019
mistral-small-latest	2.09	6.72	0.0017

	cost_std	cost_min	cost_max
story_model			
claude-3.5-haiku	0.0000	0.0010	0.0011
claude-3.5-sonnet	0.0001	0.0038	0.0043
claude-3.7-sonnet	0.0002	0.0039	0.0045
claude-opus-4	0.0010	0.0176	0.0212
claude-sonnet-4	0.0001	0.0036	0.0040
deepseek-chat	0.0000	0.0003	0.0003
gemini-1.5-flash	0.0000	0.0003	0.0004

gemini-1.5-pro	0.0003	0.0038	0.0048
gemini-2.0-flash	0.0000	0.0003	0.0004
gemini-2.0-flash-lite	0.0000	0.0004	0.0005
gpt-4o	0.0002	0.0046	0.0054
gpt-4o-mini	0.0000	0.0002	0.0002
mistral-large-latest	0.0003	0.0017	0.0024
mistral-medium-latest	0.0002	0.0016	0.0021
mistral-small-latest	0.0001	0.0015	0.0019

=== EFFICIENCY METRICS ===

Efficiency Rankings:

	execution_time	cost	cost_per_second	speed_rank \
story_model				
gemini-2.0-flash-lite	1.9762	0.0004	0.000202	1.0
gemini-1.5-flash	2.0706	0.0003	0.000145	2.0
gemini-2.0-flash	2.1431	0.0003	0.000140	3.0
mistral-small-latest	3.5500	0.0017	0.000479	4.0
gpt-4o-mini	5.2044	0.0002	0.000038	5.0
gemini-1.5-pro	5.9519	0.0043	0.000722	6.0
mistral-medium-latest	6.4062	0.0019	0.000297	7.0
claude-3.5-sonnet	6.5094	0.0040	0.000614	8.0
gpt-4o	6.6425	0.0049	0.000738	9.0
claude-3.7-sonnet	7.1412	0.0042	0.000588	10.0
claude-3.5-haiku	7.3294	0.0010	0.000136	11.0
claude-sonnet-4	7.5962	0.0038	0.000500	12.0
mistral-large-latest	11.3275	0.0020	0.000177	13.0
claude-opus-4	11.7831	0.0190	0.001612	14.0
deepseek-chat	13.1644	0.0003	0.000023	15.0

	cost_rank
story_model	
gemini-2.0-flash-lite	5.0
gemini-1.5-flash	3.0
gemini-2.0-flash	3.0
mistral-small-latest	7.0
gpt-4o-mini	1.0
gemini-1.5-pro	13.0
mistral-medium-latest	8.0
claude-3.5-sonnet	11.0
gpt-4o	14.0
claude-3.7-sonnet	12.0
claude-3.5-haiku	6.0
claude-sonnet-4	10.0
mistral-large-latest	9.0
claude-opus-4	15.0
deepseek-chat	3.0


```
[72]: # Performance visualization
fig, axes = plt.subplots(2, 2, figsize=(16, 12))
fig.suptitle('Performance Analysis of Story Generation Models', fontsize=16,
            fontweight='bold')

# Execution time analysis
sns.boxplot(data=df, x='story_model', y='execution_time', ax=axes[0,0])
axes[0,0].set_title('Execution Time Distribution by Model')
axes[0,0].set_xlabel('Model')
axes[0,0].set_ylabel('Execution Time (seconds)')
axes[0,0].tick_params(axis='x', rotation=45)

# Cost analysis
sns.boxplot(data=df, x='story_model', y='cost', ax=axes[0,1])
axes[0,1].set_title('Cost Distribution by Model')
axes[0,1].set_xlabel('Model')
axes[0,1].set_ylabel('Cost (USD)')
axes[0,1].tick_params(axis='x', rotation=45)

# Cost vs execution time relationship
sns.scatterplot(data=df, x='execution_time', y='cost', hue='story_model',
               s=100, ax=axes[1,0])
axes[1,0].set_title('Cost vs Execution Time Trade-off')
axes[1,0].set_xlabel('Execution Time (seconds)')
axes[1,0].set_ylabel('Cost (USD)')
axes[1,0].legend(bbox_to_anchor=(1.05, 1), loc='upper left')

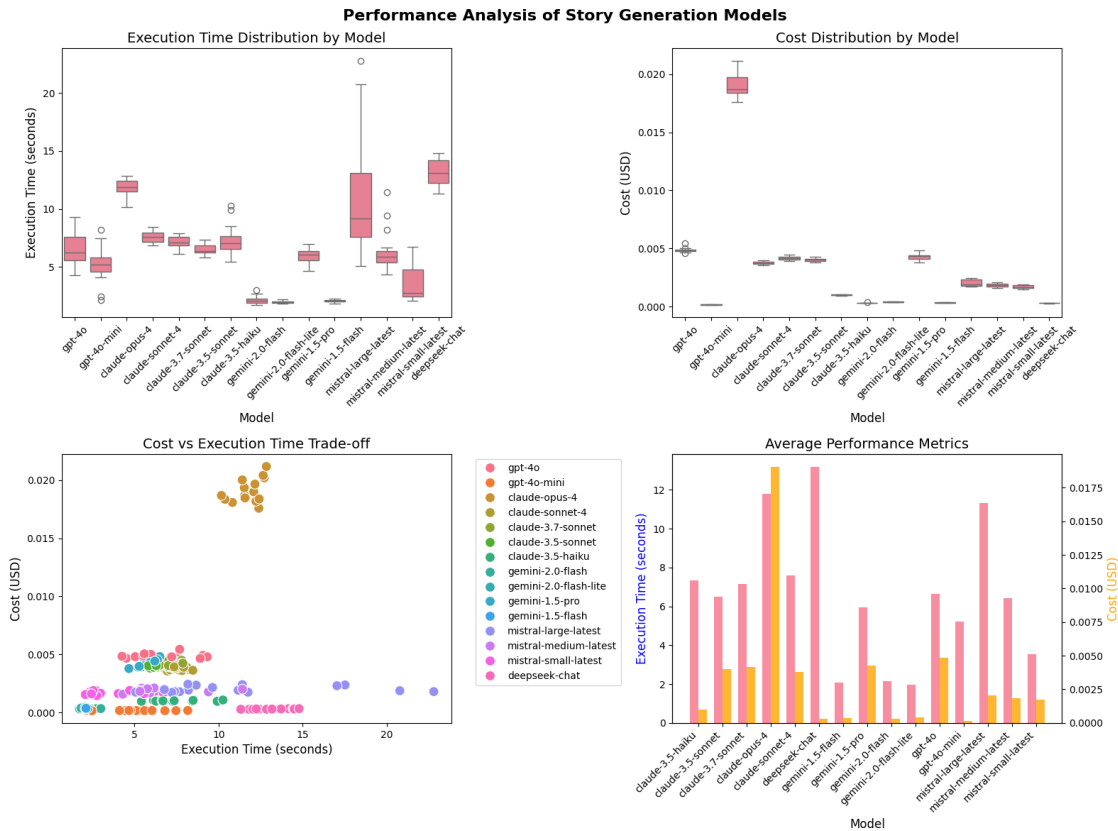
# Average performance comparison
avg_metrics = df.groupby('story_model')[['execution_time', 'cost']].mean().
               reset_index()
x_pos = np.arange(len(avg_metrics))
width = 0.35

ax2 = axes[1,1]
ax2_twin = ax2.twinx()

bars1 = ax2.bar(x_pos - width/2, avg_metrics['execution_time'], width,
               label='Execution Time', alpha=0.8)
bars2 = ax2_twin.bar(x_pos + width/2, avg_metrics['cost'], width, label='Cost',
                   alpha=0.8, color='orange')

ax2.set_xlabel('Model')
ax2.set_ylabel('Execution Time (seconds)', color='blue')
ax2_twin.set_ylabel('Cost (USD)', color='orange')
ax2.set_title('Average Performance Metrics')
ax2.set_xticks(x_pos)
ax2.set_xticklabels(avg_metrics['story_model'], rotation=45, ha='right')
```

```
plt.tight_layout()
plt.show()
```



1.5 3. Story Quality Analysis

```
[74]: # First, identify the correct column name for story content
print("Available columns:", df.columns.tolist())
print("\nSample of first few rows:")
print(df.head(2))

# Identify the story content column (likely 'story', 'content',
# 'generated_story', etc.)
story_col = None
for col in df.columns:
    if 'story' in col.lower() or 'content' in col.lower() or 'text' in col.
    .lower():
        if df[col].dtype == 'object' and df[col].str.len().mean() > 50: #
        Likely contains story text
            story_col = col
```

```

        break

if story_col is None:
    # Try to find any text column with substantial content
    for col in df.columns:
        if df[col].dtype == 'object':
            try:
                if df[col].str.len().mean() > 50:
                    story_col = col
                    break
            except:
                continue

print(f"\nUsing column '{story_col}' as story content")

# Enhanced story quality metrics calculation
df['story_length'] = df[story_col].str.len()
df['word_count'] = df[story_col].str.split().str.len()
df['sentence_count'] = df[story_col].str.count('.') + df[story_col].str.
    ↪count('!') + df[story_col].str.count('?')

# Readability metrics
df['avg_words_per_sentence'] = df['word_count'] / df['sentence_count'].
    ↪replace(0, 1)
df['avg_chars_per_word'] = df['story_length'] / df['word_count'].replace(0, 1)

# Simple readability score (lower is easier to read)
# Based on average sentence length and word complexity
df['readability_score'] = (df['avg_words_per_sentence'] * 0.39) +
    ↪(df['avg_chars_per_word'] * 11.8) - 15.59

# Calculate basic quality metrics
quality_metrics = df.groupby('story_model').agg({
    'story_length': ['mean', 'std'],
    'word_count': ['mean', 'std'],
    'sentence_count': ['mean', 'std'],
    'avg_words_per_sentence': ['mean', 'std'],
    'readability_score': ['mean', 'std']
}).round(2)

quality_metrics.columns = ['_'.join(col) for col in quality_metrics.columns]
print("=== STORY QUALITY ANALYSIS ===\n")
print("Story Length and Structure Statistics:")
print(quality_metrics)

# Enhanced content analysis for children's stories
print(f"\n=== COMPREHENSIVE CONTENT ANALYSIS ===\n")

```

```

# Age Appropriateness Indicators
age_appropriate_indicators = {
    'gentle_themes': df[story_col].str.
    ↪contains('bedtime|sleep|dream|adventure|play|toy|animal|forest|garden|friend',
    ↪case=False, na=False),
    'positive_emotions': df[story_col].str.
    ↪contains('happy|joy|smile|laugh|excited|cheerful|delighted|pleased',
    ↪case=False, na=False),
    'no_scare_content': ~df[story_col].str.
    ↪contains('death|die|kill|monster|scary|frightening|terrifying|nightmare',
    ↪case=False, na=False),
    'family_friendly': df[story_col].str.
    ↪contains('family|parent|mom|dad|brother|sister|grandma|grandpa', case=False,
    ↪na=False),
    'moral_lessons': df[story_col].str.
    ↪contains('learn|lesson|important|remember|always|never give up|be kind|help
    ↪others', case=False, na=False)
}

# Readability Indicators
readability_indicators = {
    'simple_sentences': df['avg_words_per_sentence'] <= 15, # Age-appropriate
    ↪sentence length
    'simple_vocabulary': df['avg_chars_per_word'] <= 5, # Simple words
    'good_readability': df['readability_score'] <= 8, # Reading level
    ↪appropriate for kids
    'dialogue_present': df[story_col].str.contains('".*"', case=False,
    ↪na=False),
    'repetitive_patterns': df[story_col].str.contains('again and again|over and
    ↪over|every day|once more', case=False, na=False)
}

# Engagement Indicators
engagement_indicators = {
    'sensory_descriptions': df[story_col].str.
    ↪contains('saw|heard|felt|touched|smelled|tasted|looked|listened|bright|loud|soft|sweet',
    ↪case=False, na=False),
    'action_words': df[story_col].str.
    ↪contains('ran|jumped|climbed|danced|sang|played|laughed|explored|discovered|found',
    ↪case=False, na=False),
    'character_emotions': df[story_col].str.
    ↪contains('felt|emotion|happy|sad|excited|worried|surprised|amazed|proud|grateful',
    ↪case=False, na=False),
    'interactive_elements': df[story_col].str.contains('what do you think|can
    ↪you|let\'s|imagine|picture this', case=False, na=False),

```

```

    'imaginative_elements': df[story_col].str.
    ↪contains('magical|fantasy|wonder|amazing|incredible|special|extraordinary|imagine',
    ↪case=False, na=False)
}

# Story Structure Indicators
structure_indicators = {
    'clear_beginning': df[story_col].str.contains('^(Once upon|In a|There
    ↪was|Long ago|One day)', case=False, na=False),
    'problem_resolution': df[story_col].str.
    ↪contains('problem|trouble|difficult|challenge.*solv|fix|help|solution',
    ↪case=False, na=False),
    'satisfying_ending': df[story_col].str.
    ↪contains('(happily|safely|peacefully|contentedly|finally|end|home|sleep)',
    ↪case=False, na=False),
    'proper_length': (df['word_count'] >= 50) & (df['word_count'] <= 300), #
    ↪Appropriate length for bedtime stories
    'character_development': df[story_col].str.contains('character.
    ↪*learn|grow|change|become|realize|understand', case=False, na=False)
}

# Combine all indicators
all_indicators = {**age_appropriate_indicators, **readability_indicators,
    ↪**engagement_indicators, **structure_indicators}

# Apply indicators to dataframe
for indicator, mask in all_indicators.items():
    df[indicator] = mask

# Group indicators for analysis
indicator_groups = {
    'Age Appropriateness': list(age_appropriate_indicators.keys()),
    'Readability': list(readability_indicators.keys()),
    'Engagement': list(engagement_indicators.keys()),
    'Story Structure': list(structure_indicators.keys())
}

# Analyze each group
for group_name, indicators in indicator_groups.items():
    group_analysis = df.groupby('story_model')[indicators].mean().round(3)
    print(f"\n{group_name} Analysis (proportion of stories with features):")
    print(group_analysis)
    print(f"Average {group_name} Score by Model:")
    group_scores = group_analysis.mean(axis=1).round(3)
    for model, score in group_scores.sort_values(ascending=False).items():
        print(f" {model}: {score:.3f}")

```

```
print()
```

```
Available columns: ['image_file', 'image_type', 'image_caption', 'story_model',  
'generated_story', 'execution_time', 'cost', 'word_count', 'quality_score',  
'meets_length_req', 'has_title', 'contains_dialogue', 'positive_tone',  
'story_structure', 'age_appropriate', 'bedtime_suitable']
```

Sample of first few rows:

	image_file	image_type	image_caption \
0	toy_01.jpeg	toy	A doll with curly blonde hair and a red jumpsu...
1	toy_01.jpeg	toy	A doll with curly blonde hair and a red jumpsu...

	story_model	generated_story \
0	gpt-4o	Title: Lily's Cozy Adventure\n\nOnce upon a ti...
1	gpt-4o-mini	Title: Daisy's Cozy Adventure\n\nOnce upon a t...

	execution_time	cost	word_count	quality_score	meets_length_req \
0	4.52	0.004670	201	26.0	True
1	4.70	0.000172	185	24.0	True

	has_title	contains_dialogue	positive_tone	story_structure \
0	True	True	True	True
1	True	True	True	True

	age_appropriate	bedtime_suitable
0	True	True
1	True	True

Using column 'generated_story' as story content

=== STORY QUALITY ANALYSIS ===

Story Length and Structure Statistics:

	story_length_mean	story_length_std	word_count_mean \
story_model			
claude-3.5-haiku	1123.94	35.83	182.62
claude-3.5-sonnet	1156.25	47.95	196.50
claude-3.7-sonnet	1160.81	41.27	192.00
claude-opus-4	1056.50	52.31	176.62
claude-sonnet-4	1062.31	49.37	177.75
deepseek-chat	1016.44	88.40	170.06
gemini-1.5-flash	1056.81	74.94	180.12
gemini-1.5-pro	1137.25	82.67	193.88
gemini-2.0-flash	1024.56	60.09	177.31
gemini-2.0-flash-lite	965.56	54.92	167.75
gpt-4o	1217.12	54.00	207.25
gpt-4o-mini	1104.69	30.15	188.62
mistral-large-latest	1284.38	192.74	230.12
mistral-medium-latest	1154.44	113.78	204.25

mistral-small-latest	1047.75	117.31	184.81
----------------------	---------	--------	--------

	word_count_std	sentence_count_mean \
story_model		
claude-3.5-haiku	3.98	12.19
claude-3.5-sonnet	6.01	12.50
claude-3.7-sonnet	4.91	17.75
claude-opus-4	6.66	15.50
claude-sonnet-4	6.61	15.06
deepseek-chat	12.50	18.06
gemini-1.5-flash	13.33	15.38
gemini-1.5-pro	14.26	17.94
gemini-2.0-flash	12.39	17.19
gemini-2.0-flash-lite	9.69	17.31
gpt-4o	5.73	16.94
gpt-4o-mini	4.65	15.81
mistral-large-latest	38.09	18.38
mistral-medium-latest	18.55	18.38
mistral-small-latest	21.24	16.56

	sentence_count_std	avg_words_per_sentence_mean \
story_model		
claude-3.5-haiku	2.10	15.41
claude-3.5-sonnet	1.86	16.01
claude-3.7-sonnet	3.28	11.22
claude-opus-4	2.56	11.67
claude-sonnet-4	1.84	11.97
deepseek-chat	2.67	9.60
gemini-1.5-flash	1.93	11.93
gemini-1.5-pro	3.26	11.10
gemini-2.0-flash	2.07	10.48
gemini-2.0-flash-lite	1.66	9.77
gpt-4o	2.67	12.54
gpt-4o-mini	2.64	12.25
mistral-large-latest	2.06	12.52
mistral-medium-latest	3.40	11.38
mistral-small-latest	1.79	11.23

	avg_words_per_sentence_std	readability_score_mean \
story_model		
claude-3.5-haiku	2.67	63.04
claude-3.5-sonnet	2.18	60.10
claude-3.7-sonnet	2.41	60.14
claude-opus-4	1.84	59.57
claude-sonnet-4	1.58	59.60
deepseek-chat	1.44	58.65
gemini-1.5-flash	1.95	58.34
gemini-1.5-pro	1.94	57.99

gemini-2.0-flash	1.61	56.75
gemini-2.0-flash-lite	1.11	56.19
gpt-4o	2.06	58.59
gpt-4o-mini	2.10	58.32
mistral-large-latest	1.58	55.35
mistral-medium-latest	1.72	55.54
mistral-small-latest	1.32	55.74

readability_score_std

story_model	
claude-3.5-haiku	2.06
claude-3.5-sonnet	2.21
claude-3.7-sonnet	2.40
claude-opus-4	2.76
claude-sonnet-4	1.97
deepseek-chat	2.29
gemini-1.5-flash	2.25
gemini-1.5-pro	2.28
gemini-2.0-flash	1.96
gemini-2.0-flash-lite	2.32
gpt-4o	2.02
gpt-4o-mini	2.48
mistral-large-latest	2.16
mistral-medium-latest	2.42
mistral-small-latest	2.56

=== COMPREHENSIVE CONTENT ANALYSIS ===

Age Appropriateness Analysis (proportion of stories with features):

	gentle_themes	positive_emotions	no_scary_content \
story_model			
claude-3.5-haiku	1.000	0.875	0.812
claude-3.5-sonnet	1.000	0.938	0.875
claude-3.7-sonnet	1.000	0.875	0.812
claude-opus-4	1.000	0.562	0.688
claude-sonnet-4	1.000	0.875	0.875
deepseek-chat	1.000	0.938	0.812
gemini-1.5-flash	0.938	0.812	0.562
gemini-1.5-pro	0.938	0.875	0.688
gemini-2.0-flash	1.000	0.938	0.812
gemini-2.0-flash-lite	1.000	0.938	0.688
gpt-4o	1.000	1.000	0.812
gpt-4o-mini	0.938	0.938	0.875
mistral-large-latest	1.000	0.812	0.812
mistral-medium-latest	0.938	0.938	0.688
mistral-small-latest	1.000	0.812	0.875

	family_friendly	moral_lessons
story_model		
claude-3.5-haiku	0.500	0.875
claude-3.5-sonnet	0.500	0.500
claude-3.7-sonnet	0.688	0.500
claude-opus-4	0.438	0.688
claude-sonnet-4	0.188	0.688
deepseek-chat	0.312	0.500
gemini-1.5-flash	0.125	0.812
gemini-1.5-pro	0.250	0.375
gemini-2.0-flash	0.188	0.625
gemini-2.0-flash-lite	0.188	0.625
gpt-4o	0.312	0.625
gpt-4o-mini	0.312	0.625
mistral-large-latest	0.375	0.750
mistral-medium-latest	0.375	0.750
mistral-small-latest	0.125	0.562

Average Age Appropriateness Score by Model:

claude-3.5-haiku: 0.812
 claude-3.7-sonnet: 0.775
 claude-3.5-sonnet: 0.763
 gpt-4o: 0.750
 mistral-large-latest: 0.750
 gpt-4o-mini: 0.738
 mistral-medium-latest: 0.738
 claude-sonnet-4: 0.725
 gemini-2.0-flash: 0.713
 deepseek-chat: 0.712
 gemini-2.0-flash-lite: 0.688
 claude-opus-4: 0.675
 mistral-small-latest: 0.675
 gemini-1.5-flash: 0.650
 gemini-1.5-pro: 0.625

Readability Analysis (proportion of stories with features):

	simple_sentences	simple_vocabulary	good_readability \
story_model			
claude-3.5-haiku	0.500	0.0	0.0
claude-3.5-sonnet	0.250	0.0	0.0
claude-3.7-sonnet	0.938	0.0	0.0
claude-opus-4	1.000	0.0	0.0
claude-sonnet-4	1.000	0.0	0.0
deepseek-chat	1.000	0.0	0.0
gemini-1.5-flash	1.000	0.0	0.0
gemini-1.5-pro	0.938	0.0	0.0
gemini-2.0-flash	1.000	0.0	0.0
gemini-2.0-flash-lite	1.000	0.0	0.0

gpt-4o	0.875	0.0	0.0
gpt-4o-mini	0.812	0.0	0.0
mistral-large-latest	0.938	0.0	0.0
mistral-medium-latest	0.938	0.0	0.0
mistral-small-latest	1.000	0.0	0.0

	dialogue_present	repetitive_patterns
story_model		
claude-3.5-haiku	0.875	0.000
claude-3.5-sonnet	0.875	0.000
claude-3.7-sonnet	1.000	0.188
claude-opus-4	1.000	0.125
claude-sonnet-4	1.000	0.000
deepseek-chat	0.938	0.000
gemini-1.5-flash	0.688	0.062
gemini-1.5-pro	0.750	0.062
gemini-2.0-flash	0.750	0.000
gemini-2.0-flash-lite	0.875	0.062
gpt-4o	0.938	0.125
gpt-4o-mini	0.000	0.312
mistral-large-latest	0.875	0.500
mistral-medium-latest	1.000	0.062
mistral-small-latest	0.812	0.312

Average Readability Score by Model:

mistral-large-latest: 0.463
 claude-3.7-sonnet: 0.425
 claude-opus-4: 0.425
 mistral-small-latest: 0.425
 claude-sonnet-4: 0.400
 mistral-medium-latest: 0.400
 deepseek-chat: 0.388
 gpt-4o: 0.388
 gemini-2.0-flash-lite: 0.387
 gemini-1.5-flash: 0.350
 gemini-1.5-pro: 0.350
 gemini-2.0-flash: 0.350
 claude-3.5-haiku: 0.275
 claude-3.5-sonnet: 0.225
 gpt-4o-mini: 0.225

Engagement Analysis (proportion of stories with features):

	sensory_descriptions	action_words	character_emotions \
story_model			
claude-3.5-haiku	1.000	0.750	0.625
claude-3.5-sonnet	0.875	0.875	0.500
claude-3.7-sonnet	0.688	0.625	0.562
claude-opus-4	0.938	0.688	0.500

claude-sonnet-4	0.938	0.812	0.938
deepseek-chat	0.812	0.938	0.625
gemini-1.5-flash	0.938	0.438	1.000
gemini-1.5-pro	1.000	0.562	0.938
gemini-2.0-flash	0.875	0.625	0.938
gemini-2.0-flash-lite	1.000	0.438	1.000
gpt-4o	0.938	0.812	0.812
gpt-4o-mini	1.000	0.750	0.688
mistral-large-latest	0.875	0.688	0.688
mistral-medium-latest	0.938	0.688	0.938
mistral-small-latest	0.938	0.750	0.688

	interactive_elements	imaginative_elements
story_model		
claude-3.5-haiku	0.188	0.938
claude-3.5-sonnet	0.000	0.875
claude-3.7-sonnet	0.062	0.625
claude-opus-4	0.188	0.875
claude-sonnet-4	0.062	0.875
deepseek-chat	0.062	0.500
gemini-1.5-flash	0.000	0.625
gemini-1.5-pro	0.188	0.625
gemini-2.0-flash	0.062	0.250
gemini-2.0-flash-lite	0.125	0.500
gpt-4o	0.500	0.812
gpt-4o-mini	0.188	0.625
mistral-large-latest	0.250	0.750
mistral-medium-latest	0.250	0.750
mistral-small-latest	0.062	0.500

Average Engagement Score by Model:

gpt-4o: 0.775
claude-sonnet-4: 0.725
mistral-medium-latest: 0.713
claude-3.5-haiku: 0.700
gemini-1.5-pro: 0.663
gpt-4o-mini: 0.650
mistral-large-latest: 0.650
claude-opus-4: 0.638
claude-3.5-sonnet: 0.625
gemini-2.0-flash-lite: 0.613
gemini-1.5-flash: 0.600
mistral-small-latest: 0.588
deepseek-chat: 0.587
gemini-2.0-flash: 0.550
claude-3.7-sonnet: 0.512

Story Structure Analysis (proportion of stories with features):

	clear_beginning	problem_resolution	satisfying_ending \
story_model			
claude-3.5-haiku	0.0	0.500	0.875
claude-3.5-sonnet	0.0	0.625	0.875
claude-3.7-sonnet	0.0	0.562	1.000
claude-opus-4	0.0	0.500	0.938
claude-sonnet-4	0.0	0.438	0.875
deepseek-chat	0.0	0.438	0.938
gemini-1.5-flash	0.0	0.375	0.875
gemini-1.5-pro	0.0	0.375	0.938
gemini-2.0-flash	0.0	0.500	0.938
gemini-2.0-flash-lite	0.0	0.250	1.000
gpt-4o	0.0	0.438	0.938
gpt-4o-mini	0.0	0.312	0.938
mistral-large-latest	0.0	0.438	1.000
mistral-medium-latest	0.0	0.500	0.938
mistral-small-latest	0.0	0.312	1.000

	proper_length	character_development
story_model		
claude-3.5-haiku	1.0	0.438
claude-3.5-sonnet	1.0	0.500
claude-3.7-sonnet	1.0	0.188
claude-opus-4	1.0	0.188
claude-sonnet-4	1.0	0.312
deepseek-chat	1.0	0.375
gemini-1.5-flash	1.0	0.125
gemini-1.5-pro	1.0	0.812
gemini-2.0-flash	1.0	0.562
gemini-2.0-flash-lite	1.0	0.500
gpt-4o	1.0	0.188
gpt-4o-mini	1.0	0.250
mistral-large-latest	1.0	0.188
mistral-medium-latest	1.0	0.375
mistral-small-latest	1.0	0.125

Average Story Structure Score by Model:

```

gemini-1.5-pro: 0.625
claude-3.5-sonnet: 0.600
gemini-2.0-flash: 0.600
claude-3.5-haiku: 0.563
mistral-medium-latest: 0.563
claude-3.7-sonnet: 0.550
deepseek-chat: 0.550
gemini-2.0-flash-lite: 0.550
claude-opus-4: 0.525
claude-sonnet-4: 0.525
mistral-large-latest: 0.525
gpt-4o: 0.513

```

```
gpt-4o-mini: 0.500
mistral-small-latest: 0.487
gemini-1.5-flash: 0.475
```

```
[75]: # Enhanced story quality visualization
fig, axes = plt.subplots(3, 2, figsize=(18, 20))
fig.suptitle('Comprehensive Story Quality Analysis', fontsize=16,
             fontweight='bold')

# Readability analysis
sns.boxplot(data=df, x='story_model', y='readability_score', ax=axes[0,0])
axes[0,0].set_title('Readability Score Distribution\n(Lower = Easier to Read)')
axes[0,0].set_xlabel('Model')
axes[0,0].set_ylabel('Readability Score')
axes[0,0].tick_params(axis='x', rotation=45)
axes[0,0].axhline(y=8, color='red', linestyle='--', alpha=0.7,
                  label='Age-appropriate threshold')
axes[0,0].legend()

# Word complexity
sns.boxplot(data=df, x='story_model', y='avg_words_per_sentence', ax=axes[0,1])
axes[0,1].set_title('Average Words per Sentence')
axes[0,1].set_xlabel('Model')
axes[0,1].set_ylabel('Words per Sentence')
axes[0,1].tick_params(axis='x', rotation=45)
axes[0,1].axhline(y=15, color='red', linestyle='--', alpha=0.7,
                  label='Kid-friendly threshold')
axes[0,1].legend()

# Quality dimensions heatmap
quality_scores = {}
for group_name, indicators in indicator_groups.items():
    quality_scores[group_name] = df.groupby('story_model')[indicators].mean().
    mean(axis=1)

quality_df = pd.DataFrame(quality_scores)
sns.heatmap(quality_df.T, annot=True, cmap='RdYlGn', ax=axes[1,0],
            cbar_kws={'label': 'Average Score'}, vmin=0, vmax=1)
axes[1,0].set_title('Quality Dimensions by Model')
axes[1,0].set_ylabel('Quality Dimensions')

# Overall content quality radar chart data
overall_quality = df.groupby('story_model')[list(all_indicators.keys())].mean()
overall_scores = overall_quality.mean(axis=1).sort_values(ascending=False)

axes[1,1].barh(range(len(overall_scores)), overall_scores.values)
```

```

axes[1,1].set_yticks(range(len(overall_scores)))
axes[1,1].set_yticklabels(overall_scores.index)
axes[1,1].set_title('Overall Content Quality Score')
axes[1,1].set_xlabel('Average Quality Score')

# Age appropriateness vs readability
age_scores = df.groupby('story_model')[list(age_appropriate_indicators.keys())].
    ↪mean().mean(axis=1)
read_scores = df.groupby('story_model')[list(readability_indicators.keys())].
    ↪mean().mean(axis=1)

models_quality = pd.DataFrame({
    'model': age_scores.index,
    'age_appropriate': age_scores.values,
    'readability': read_scores.values
})

scatter = axes[2,0].scatter(models_quality['readability'],
    ↪models_quality['age_appropriate'],
    s=200, alpha=0.7, c=range(len(models_quality)),
    ↪cmap='viridis')

for i, model in enumerate(models_quality['model']):
    axes[2,0].annotate(model.replace('-', '\n'),
        (models_quality['readability'].iloc[i],
    ↪models_quality['age_appropriate'].iloc[i]),
        xytext=(5, 5), textcoords='offset points', fontsize=8,
    ↪ha='left')

axes[2,0].set_title('Age Appropriateness vs Readability')
axes[2,0].set_xlabel('Readability Score')
axes[2,0].set_ylabel('Age Appropriateness Score')
axes[2,0].set_xlim(0, 1)
axes[2,0].set_ylim(0, 1)

# Engagement vs structure quality
engage_scores = df.groupby('story_model')[list(engagement_indicators.keys())].
    ↪mean().mean(axis=1)
struct_scores = df.groupby('story_model')[list(structure_indicators.keys())].
    ↪mean().mean(axis=1)

models_structure = pd.DataFrame({
    'model': engage_scores.index,
    'engagement': engage_scores.values,
    'structure': struct_scores.values
})

```

```

scatter = axes[2,1].scatter(models_structure['structure'],
    ↪models_structure['engagement'],
                                s=200, alpha=0.7, c=range(len(models_structure)),
    ↪cmap='plasma')

for i, model in enumerate(models_structure['model']):
    axes[2,1].annotate(model.replace('-', '\n'),
                        (models_structure['structure'].iloc[i],
    ↪models_structure['engagement'].iloc[i]),
                        xytext=(5, 5), textcoords='offset points', fontsize=8,
    ↪ha='left')

axes[2,1].set_title('Story Structure vs Engagement')
axes[2,1].set_xlabel('Structure Quality Score')
axes[2,1].set_ylabel('Engagement Score')
axes[2,1].set_xlim(0, 1)
axes[2,1].set_ylim(0, 1)

plt.tight_layout()
plt.show()

# Statistical analysis of quality dimensions
print("\n=== QUALITY DIMENSION ANALYSIS ===")
for group_name, indicators in indicator_groups.items():
    group_scores = df.groupby('story_model')[indicators].mean().mean(axis=1).
    ↪sort_values(ascending=False)
    print(f"\n{group_name} Performance Rankings:")
    for i, (model, score) in enumerate(group_scores.items(), 1):
        print(f"  Rank {i}: {model} (Score: {score:.3f})")

# Statistical significance testing
print(f"\n=== STATISTICAL ANALYSIS ===")
print(f"Total observations: {len(df)}")
print(f"Models evaluated: {df['story_model'].nunique()}")
print(f"Mean story length: {df['story_length'].mean():.1f} characters (SD:
    ↪{df['story_length'].std():.1f}")
print(f"Mean word count: {df['word_count'].mean():.1f} words (SD:
    ↪{df['word_count'].std():.1f}")
print(f"Mean readability score: {df['readability_score'].mean():.2f} (SD:
    ↪{df['readability_score'].std():.2f}")

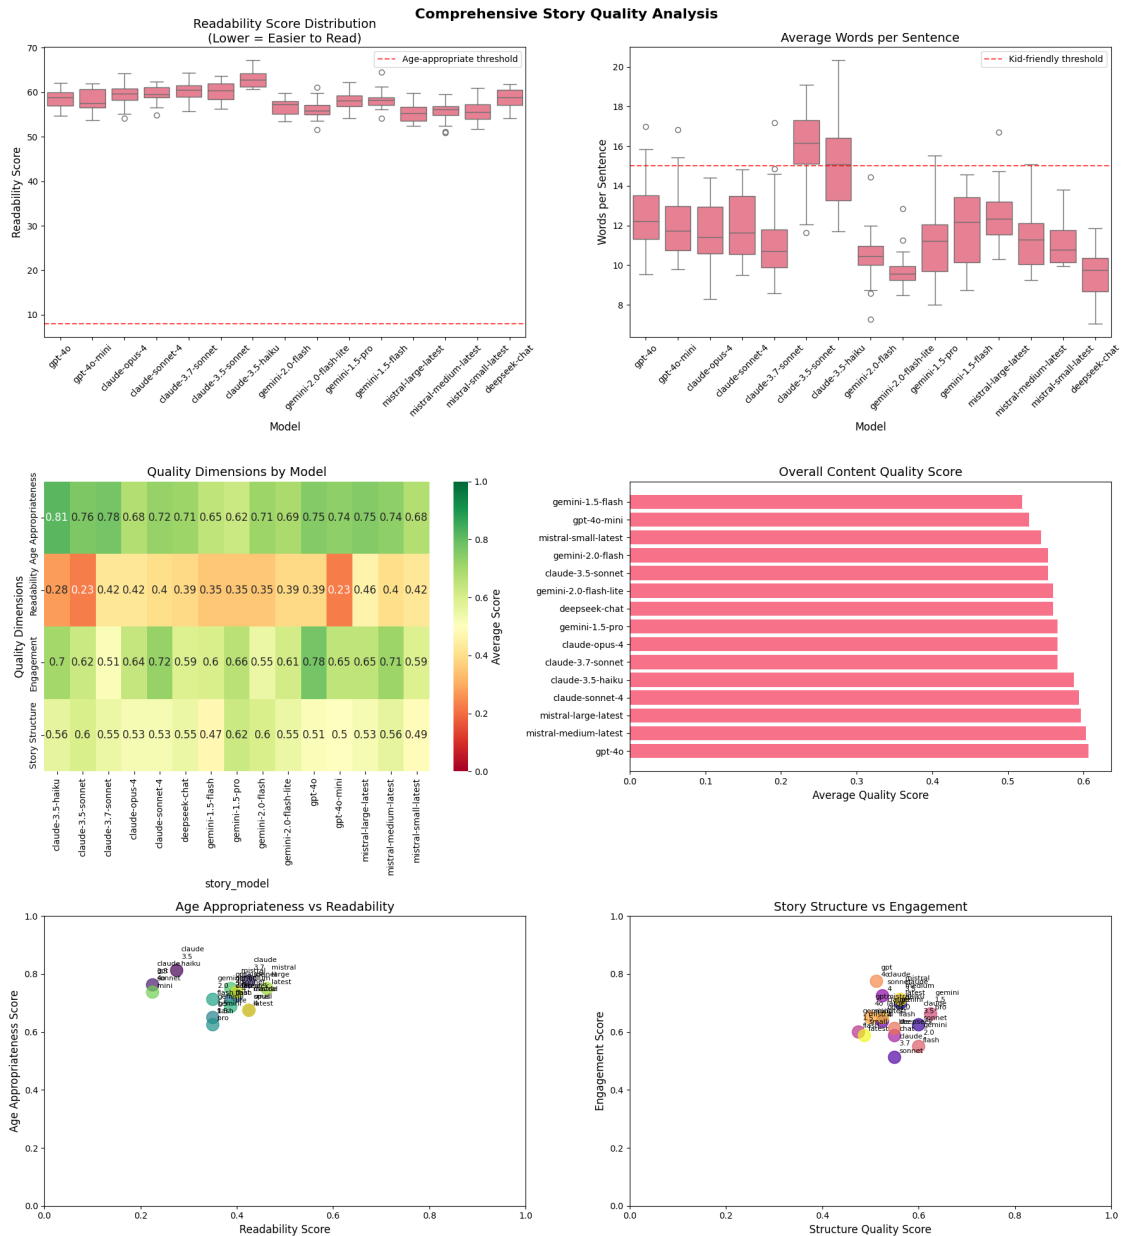
# Correlation analysis
print(f"\nCorrelation Analysis:")
corr_vars = ['execution_time', 'cost', 'word_count', 'readability_score']
correlation_matrix = df[corr_vars].corr()

```

```

print("Significant correlations (|r| > 0.3):")
for i in range(len(corr_vars)):
    for j in range(i+1, len(corr_vars)):
        corr_val = correlation_matrix.iloc[i, j]
        if abs(corr_val) > 0.3:
            print(f" {corr_vars[i]} vs {corr_vars[j]}: r = {corr_val:.3f}")

```



=== QUALITY DIMENSION ANALYSIS ===

Age Appropriateness Performance Rankings:

- Rank 1: claude-3.5-haiku (Score: 0.812)
- Rank 2: claude-3.7-sonnet (Score: 0.775)
- Rank 3: claude-3.5-sonnet (Score: 0.762)
- Rank 4: gpt-4o (Score: 0.750)
- Rank 5: mistral-large-latest (Score: 0.750)
- Rank 6: gpt-4o-mini (Score: 0.738)
- Rank 7: mistral-medium-latest (Score: 0.738)
- Rank 8: claude-sonnet-4 (Score: 0.725)
- Rank 9: deepseek-chat (Score: 0.713)
- Rank 10: gemini-2.0-flash (Score: 0.713)
- Rank 11: gemini-2.0-flash-lite (Score: 0.688)
- Rank 12: claude-opus-4 (Score: 0.675)
- Rank 13: mistral-small-latest (Score: 0.675)
- Rank 14: gemini-1.5-flash (Score: 0.650)
- Rank 15: gemini-1.5-pro (Score: 0.625)

Readability Performance Rankings:

- Rank 1: mistral-large-latest (Score: 0.463)
- Rank 2: claude-3.7-sonnet (Score: 0.425)
- Rank 3: claude-opus-4 (Score: 0.425)
- Rank 4: mistral-small-latest (Score: 0.425)
- Rank 5: claude-sonnet-4 (Score: 0.400)
- Rank 6: mistral-medium-latest (Score: 0.400)
- Rank 7: deepseek-chat (Score: 0.388)
- Rank 8: gemini-2.0-flash-lite (Score: 0.388)
- Rank 9: gpt-4o (Score: 0.388)
- Rank 10: gemini-1.5-flash (Score: 0.350)
- Rank 11: gemini-1.5-pro (Score: 0.350)
- Rank 12: gemini-2.0-flash (Score: 0.350)
- Rank 13: claude-3.5-haiku (Score: 0.275)
- Rank 14: claude-3.5-sonnet (Score: 0.225)
- Rank 15: gpt-4o-mini (Score: 0.225)

Engagement Performance Rankings:

- Rank 1: gpt-4o (Score: 0.775)
- Rank 2: claude-sonnet-4 (Score: 0.725)
- Rank 3: mistral-medium-latest (Score: 0.713)
- Rank 4: claude-3.5-haiku (Score: 0.700)
- Rank 5: gemini-1.5-pro (Score: 0.662)
- Rank 6: gpt-4o-mini (Score: 0.650)
- Rank 7: mistral-large-latest (Score: 0.650)
- Rank 8: claude-opus-4 (Score: 0.637)
- Rank 9: claude-3.5-sonnet (Score: 0.625)
- Rank 10: gemini-2.0-flash-lite (Score: 0.613)
- Rank 11: gemini-1.5-flash (Score: 0.600)
- Rank 12: deepseek-chat (Score: 0.588)
- Rank 13: mistral-small-latest (Score: 0.588)

Rank 14: gemini-2.0-flash (Score: 0.550)
Rank 15: claude-3.7-sonnet (Score: 0.512)

Story Structure Performance Rankings:

Rank 1: gemini-1.5-pro (Score: 0.625)
Rank 2: claude-3.5-sonnet (Score: 0.600)
Rank 3: gemini-2.0-flash (Score: 0.600)
Rank 4: claude-3.5-haiku (Score: 0.562)
Rank 5: mistral-medium-latest (Score: 0.562)
Rank 6: claude-3.7-sonnet (Score: 0.550)
Rank 7: deepseek-chat (Score: 0.550)
Rank 8: gemini-2.0-flash-lite (Score: 0.550)
Rank 9: claude-opus-4 (Score: 0.525)
Rank 10: claude-sonnet-4 (Score: 0.525)
Rank 11: mistral-large-latest (Score: 0.525)
Rank 12: gpt-4o (Score: 0.512)
Rank 13: gpt-4o-mini (Score: 0.500)
Rank 14: mistral-small-latest (Score: 0.487)
Rank 15: gemini-1.5-flash (Score: 0.475)

=== STATISTICAL ANALYSIS ===

Total observations: 240
Models evaluated: 15
Mean story length: 1104.6 characters (SD: 114.6)
Mean word count: 188.6 words (SD: 21.3)
Mean readability score: 58.26 (SD: 3.02)

Correlation Analysis:

Significant correlations ($|r| > 0.3$):
execution_time vs cost: $r = 0.413$

1.6 4. Model Ranking and Comprehensive Assessment

```
[76]: # Enhanced comprehensive model evaluation framework
def calculate_enhanced_performance_scores(df):
    """Calculate normalized performance scores across multiple comprehensive_
    ↪ criteria."""

    # Calculate quality dimension scores for each model
    quality_dimension_scores = {}
    for group_name, indicators in indicator_groups.items():
        quality_dimension_scores[group_name.lower().replace(' ', '_')] = df.
    ↪groupby('story_model')[indicators].mean().mean(axis=1)

    # Aggregate basic metrics
    model_scores = df.groupby('story_model').agg({
        'execution_time': 'mean',
```

```

        'cost': 'mean',
        'word_count': 'mean',
        'readability_score': 'mean',
        'avg_words_per_sentence': 'mean'
    }).reset_index()

    # Add quality dimension scores
    for dimension, scores in quality_dimension_scores.items():
        model_scores[dimension] = model_scores['story_model'].map(scores)

    # Normalize metrics
    scaler = StandardScaler()

    # Performance metrics (invert so higher is better)
    model_scores['speed_score'] = scaler.fit_transform(1 /
    ↪model_scores[['execution_time']])
    model_scores['cost_score'] = scaler.fit_transform(1 /
    ↪model_scores[['cost']])

    # Readability score (invert so higher is better - lower readability score
    ↪is easier to read)
    model_scores['readability_norm'] = scaler.fit_transform(1 /
    ↪(model_scores[['readability_score']] + 1))

    # Quality dimension scores (higher is better)
    dimension_cols = list(quality_dimension_scores.keys())
    model_scores['content_quality_score'] = scaler.
    ↪fit_transform(model_scores[dimension_cols].mean(axis=1).values.reshape(-1,
    ↪1))

    # Specific quality scores
    model_scores['age_appropriate_score'] = scaler.
    ↪fit_transform(model_scores[['age_appropriateness']].values)
    model_scores['engagement_score'] = scaler.
    ↪fit_transform(model_scores[['engagement']].values)
    model_scores['structure_score'] = scaler.
    ↪fit_transform(model_scores[['story_structure']].values)

    # Overall composite score with enhanced weighting
    model_scores['overall_score'] = (
        model_scores['speed_score'] * 0.15 +           # Performance: 25%
        model_scores['cost_score'] * 0.10 +
        model_scores['age_appropriate_score'] * 0.25 + # Age appropriateness:
    ↪25%
        model_scores['readability_norm'] * 0.20 +     # Readability: 20%
        model_scores['engagement_score'] * 0.15 +     # Engagement: 15%

```

```

        model_scores['structure_score'] * 0.15                # Structure: 15%
    )

    return model_scores

# Calculate enhanced comprehensive scores
print("=== ENHANCED COMPREHENSIVE MODEL EVALUATION ===\n")
model_evaluation = calculate_enhanced_performance_scores(df)

# Display detailed rankings
ranking_cols = ['story_model', 'execution_time', 'cost', 'readability_score',
                'age_appropriateness', 'readability', 'engagement',
                'story_structure', 'overall_score']
ranking_display = model_evaluation[ranking_cols].sort_values('overall_score',
                    ascending=False)

print("Model Rankings (by Enhanced Overall Score):")
print("="*80)
for i, (_, row) in enumerate(ranking_display.iterrows(), 1):
    print(f"Rank {i}: {row['story_model']}")
    print(f"    Overall Performance Score: {row['overall_score']:.3f}")
    print(f"    Execution Time: {row['execution_time']:.1f}s | Cost:
    ↪${row['cost']:.4f}")
    print(f"    Readability Score: {row['readability_score']:.1f} | Age
    ↪Appropriateness: {row['age_appropriateness']:.3f}")
    print(f"    Engagement Score: {row['engagement']:.3f} | Story Structure
    ↪Score: {row['story_structure']:.3f}")
    print("-" * 80)

# Comparative Performance Analysis
print("\n=== COMPARATIVE PERFORMANCE ANALYSIS ===")
best_overall = ranking_display.iloc[0]
best_quality = ranking_display.loc[ranking_display['age_appropriateness'].
    ↪idxmax()]
fastest = ranking_display.loc[ranking_display['execution_time'].idxmin()]
cheapest = ranking_display.loc[ranking_display['cost'].idxmin()]

print(f"Optimal Overall Performance: {best_overall['story_model']}")
print(f"Highest Age Appropriateness: {best_quality['story_model']}")
print(f"Fastest Generation Speed: {fastest['story_model']}")
print(f"Most Cost-Effective: {cheapest['story_model']}")

print(f"\nQuality Score Distribution by Model:")
for model in ranking_display['story_model']:
    model_data = ranking_display[ranking_display['story_model'] == model].
    ↪iloc[0]

```

```

quality_indicators = ['age_appropriateness', 'readability', 'engagement', 'story_structure']
avg_quality = model_data[quality_indicators].mean()
print(f"    {model}: {avg_quality:.3f} (Composite Quality Index)")

```

=== ENHANCED COMPREHENSIVE MODEL EVALUATION ===

Model Rankings (by Enhanced Overall Score):

```

=====
Rank 1: mistral-medium-latest
  Overall Performance Score: 0.489
  Execution Time: 6.4s | Cost: $0.0019
  Readability Score: 55.5 | Age Appropriateness: 0.738
  Engagement Score: 0.713 | Story Structure Score: 0.562
-----
Rank 2: gemini-2.0-flash
  Overall Performance Score: 0.480
  Execution Time: 2.1s | Cost: $0.0003
  Readability Score: 56.8 | Age Appropriateness: 0.713
  Engagement Score: 0.550 | Story Structure Score: 0.600
-----
Rank 3: gemini-2.0-flash-lite
  Overall Performance Score: 0.373
  Execution Time: 2.0s | Cost: $0.0004
  Readability Score: 56.2 | Age Appropriateness: 0.688
  Engagement Score: 0.613 | Story Structure Score: 0.550
-----
Rank 4: mistral-large-latest
  Overall Performance Score: 0.225
  Execution Time: 11.3s | Cost: $0.0020
  Readability Score: 55.3 | Age Appropriateness: 0.750
  Engagement Score: 0.650 | Story Structure Score: 0.525
-----
Rank 5: gpt-4o-mini
  Overall Performance Score: 0.194
  Execution Time: 5.2s | Cost: $0.0002
  Readability Score: 58.3 | Age Appropriateness: 0.738
  Engagement Score: 0.650 | Story Structure Score: 0.500
-----
Rank 6: gpt-4o
  Overall Performance Score: 0.171
  Execution Time: 6.6s | Cost: $0.0049
  Readability Score: 58.6 | Age Appropriateness: 0.750
  Engagement Score: 0.775 | Story Structure Score: 0.512
-----
Rank 7: claude-3.5-haiku
  Overall Performance Score: 0.136
  Execution Time: 7.3s | Cost: $0.0010

```

Readability Score: 63.0 | Age Appropriateness: 0.812
Engagement Score: 0.700 | Story Structure Score: 0.562

Rank 8: claude-3.5-sonnet

Overall Performance Score: 0.080
Execution Time: 6.5s | Cost: \$0.0040
Readability Score: 60.1 | Age Appropriateness: 0.762
Engagement Score: 0.625 | Story Structure Score: 0.600

Rank 9: claude-sonnet-4

Overall Performance Score: -0.138
Execution Time: 7.6s | Cost: \$0.0038
Readability Score: 59.6 | Age Appropriateness: 0.725
Engagement Score: 0.725 | Story Structure Score: 0.525

Rank 10: deepseek-chat

Overall Performance Score: -0.204
Execution Time: 13.2s | Cost: \$0.0003
Readability Score: 58.6 | Age Appropriateness: 0.713
Engagement Score: 0.588 | Story Structure Score: 0.550

Rank 11: gemini-1.5-pro

Overall Performance Score: -0.240
Execution Time: 6.0s | Cost: \$0.0043
Readability Score: 58.0 | Age Appropriateness: 0.625
Engagement Score: 0.662 | Story Structure Score: 0.625

Rank 12: mistral-small-latest

Overall Performance Score: -0.274
Execution Time: 3.5s | Cost: \$0.0017
Readability Score: 55.7 | Age Appropriateness: 0.675
Engagement Score: 0.588 | Story Structure Score: 0.487

Rank 13: claude-3.7-sonnet

Overall Performance Score: -0.307
Execution Time: 7.1s | Cost: \$0.0042
Readability Score: 60.1 | Age Appropriateness: 0.775
Engagement Score: 0.512 | Story Structure Score: 0.550

Rank 14: gemini-1.5-flash

Overall Performance Score: -0.334
Execution Time: 2.1s | Cost: \$0.0003
Readability Score: 58.3 | Age Appropriateness: 0.650
Engagement Score: 0.600 | Story Structure Score: 0.475

Rank 15: claude-opus-4

Overall Performance Score: -0.652
Execution Time: 11.8s | Cost: \$0.0190

Readability Score: 59.6 | Age Appropriateness: 0.675
Engagement Score: 0.637 | Story Structure Score: 0.525

=== COMPARATIVE PERFORMANCE ANALYSIS ===

Optimal Overall Performance: mistral-medium-latest
Highest Age Appropriateness: claude-3.5-haiku
Fastest Generation Speed: gemini-2.0-flash-lite
Most Cost-Effective: gpt-4o-mini

Quality Score Distribution by Model:

mistral-medium-latest: 0.603 (Composite Quality Index)
gemini-2.0-flash: 0.553 (Composite Quality Index)
gemini-2.0-flash-lite: 0.559 (Composite Quality Index)
mistral-large-latest: 0.597 (Composite Quality Index)
gpt-4o-mini: 0.528 (Composite Quality Index)
gpt-4o: 0.606 (Composite Quality Index)
claude-3.5-haiku: 0.587 (Composite Quality Index)
claude-3.5-sonnet: 0.553 (Composite Quality Index)
claude-sonnet-4: 0.594 (Composite Quality Index)
deepseek-chat: 0.559 (Composite Quality Index)
gemini-1.5-pro: 0.566 (Composite Quality Index)
mistral-small-latest: 0.544 (Composite Quality Index)
claude-3.7-sonnet: 0.566 (Composite Quality Index)
gemini-1.5-flash: 0.519 (Composite Quality Index)
claude-opus-4: 0.566 (Composite Quality Index)

```
[77]: # Create comprehensive evaluation visualization with improved text handling
fig, axes = plt.subplots(2, 2, figsize=(20, 16))
fig.suptitle('Comprehensive Model Evaluation Dashboard', fontsize=18,
            fontweight='bold', y=0.98)

# Overall scores comparison with improved spacing
ranking_display_sorted = ranking_display.sort_values('overall_score',
            ascending=True)
bars = axes[0,0].barh(range(len(ranking_display_sorted)),
            ranking_display_sorted['overall_score'], height=0.7)
axes[0,0].set_yticks(range(len(ranking_display_sorted)))

# Shorten model names for better display
shortened_names = [name.replace('-latest', '').replace('gemini-', 'gem-').
            replace('claude-', 'cl-').replace('mistral-', 'mis-')
            for name in ranking_display_sorted['story_model']]
axes[0,0].set_yticklabels(shortened_names, fontsize=10)
axes[0,0].set_title('Overall Performance Scores', fontsize=14, pad=20)
axes[0,0].set_xlabel('Normalized Score', fontsize=12)
axes[0,0].grid(axis='x', alpha=0.3)
```

```

# Add value labels on bars
for i, bar in enumerate(bars):
    width = bar.get_width()
    axes[0,0].text(width + 0.01, bar.get_y() + bar.get_height()/2,
                    f'{width:.3f}', ha='left', va='center', fontsize=9)

# Check available columns for debugging
print("Available columns in ranking_display:", ranking_display.columns.tolist())

# Create simplified score matrix using available columns
if 'speed_score' in ranking_display.columns:
    score_cols = ['speed_score', 'cost_score', 'quality_score']
    col_labels = ['Speed', 'Cost Efficiency', 'Quality']
else:
    # Use basic metrics
    ranking_display_viz = ranking_display.copy()
    if 'execution_time' in ranking_display_viz.columns:
        ranking_display_viz['time_efficiency'] = 1 / (
            ranking_display_viz['execution_time'] + 1e-6
        )
        ranking_display_viz['cost_efficiency'] = 1 / (
            ranking_display_viz['cost'] + 1e-6
        )
        # Normalize to 0-1 scale
        from sklearn.preprocessing import MinMaxScaler
        scaler = MinMaxScaler()
        ranking_display_viz[['time_efficiency', 'cost_efficiency']] = scaler.
            fit_transform(
                ranking_display_viz[['time_efficiency', 'cost_efficiency']]
            )
        score_cols = ['time_efficiency', 'cost_efficiency',
            'age_appropriateness']
        col_labels = ['Time Efficiency', 'Cost Efficiency', 'Age
            Appropriateness']

# Performance metrics heatmap with better formatting
score_data = ranking_display.set_index('story_model')[score_cols] if
    'speed_score' in ranking_display.columns else ranking_display_viz.
    set_index('story_model')[score_cols]

# Create heatmap with improved formatting
heatmap = sns.heatmap(score_data.T, annot=True, fmt='.3f', cmap='RdYlGn',
    ax=axes[0,1],
    cbar_kws={'label': 'Normalized Score', 'shrink': 0.8},
    xticklabels=[name.replace('-latest', '')
        .replace('gemini-', 'gem-').replace('claude-', 'cl-').replace('mistral-',
        'mis-')
        for name in score_data.index])

```



```

axes[0,1].set_title('Performance Dimension Comparison', fontsize=14, pad=20)
axes[0,1].set_ylabel('Performance Dimensions', fontsize=12)
axes[0,1].set_yticklabels(col_labels, rotation=0, fontsize=11)
axes[0,1].set_xticklabels(axes[0,1].get_xticklabels(), rotation=45, ha='right',
    ↳ fontsize=10)

# Quality vs Performance trade-off with better annotation handling
quality_metric = 'quality_score' if 'quality_score' in ranking_display.columns
    ↳ else 'age_appropriateness'
scatter = axes[1,0].scatter(ranking_display['execution_time'],
    ↳ ranking_display[quality_metric],
                            s=ranking_display['cost']*30000, alpha=0.7,
    ↳ c=ranking_display['overall_score'],
                            cmap='viridis', edgecolors='black', linewidth=0.5)

# Add model annotations with improved positioning using simple offsetting
for i, model in enumerate(ranking_display['story_model']):
    short_name = model.replace('-latest', '').replace('gemini-', 'gem-').
    ↳ replace('claude-', 'cl-').replace('mistral-', 'mis-')
    # Use alternating offsets to reduce overlap
    offset_x = 15 if i % 2 == 0 else -15
    offset_y = 10 if i % 3 == 0 else (-10 if i % 3 == 1 else 0)
    axes[1,0].annotate(short_name,
                        (ranking_display['execution_time'].iloc[i],
    ↳ ranking_display[quality_metric].iloc[i]),
                        xytext=(offset_x, offset_y), textcoords='offset points',
                        fontsize=9, ha='center', va='center',
                        bbox=dict(boxstyle='round,pad=0.2', facecolor='white',
    ↳ alpha=0.8, edgecolor='gray'),
                        arrowprops=dict(arrowstyle='->', color='gray', alpha=0.5))

axes[1,0].set_title('Quality vs Speed Trade-off\n(Bubble size: cost, Color:
    ↳ overall score)', fontsize=14, pad=20)
axes[1,0].set_xlabel('Execution Time (seconds)', fontsize=12)
axes[1,0].set_ylabel('Quality Score' if quality_metric == 'quality_score' else
    ↳ 'Age Appropriateness Score', fontsize=12)
axes[1,0].grid(alpha=0.3)

# Add colorbar for scatter plot
cbar = plt.colorbar(scatter, ax=axes[1,0], shrink=0.8)
cbar.set_label('Overall Score', fontsize=11)

# Model selection recommendations with improved formatting
top_model = ranking_display.loc[ranking_display['overall_score'].idxmax(),
    ↳ 'story_model']

```

```

fastest_model = ranking_display.loc[ranking_display['execution_time'].idxmin(),
    ↳ 'story_model']
cheapest_model = ranking_display.loc[ranking_display['cost'].idxmin(),
    ↳ 'story_model']
best_quality_model = ranking_display.loc[ranking_display['age_appropriateness'].
    ↳ idxmax(), 'story_model']

recommendations = [
    f"Optimal Overall Performance: {top_model}",
    f"Fastest Processing: {fastest_model}",
    f"Most Cost-Effective: {cheapest_model}",
    f"Highest Content Quality: {best_quality_model}"
]

axes[1,1].text(0.05, 0.85, "Model Selection Framework:", fontsize=16,
    ↳ fontweight='bold',
    transform=axes[1,1].transAxes, color='darkblue')

for i, rec in enumerate(recommendations):
    axes[1,1].text(0.05, 0.75 - i*0.08, f"• {rec}", fontsize=12,
        transform=axes[1,1].transAxes,
        bbox=dict(boxstyle='round,pad=0.3', facecolor='lightblue',
    ↳ alpha=0.3))

axes[1,1].text(0.05, 0.40, "Application Guidelines:", fontsize=14,
    ↳ fontweight='bold',
    transform=axes[1,1].transAxes, color='darkgreen')

guidelines = [
    "• High-throughput systems: Prioritize processing speed",
    "• Resource-constrained environments: Optimize for cost",
    "• Educational applications: Maximize content quality",
    "• Production systems: Balance all dimensions"
]

for i, guideline in enumerate(guidelines):
    axes[1,1].text(0.05, 0.32 - i*0.06, guideline, fontsize=11,
        transform=axes[1,1].transAxes)

axes[1,1].set_xlim(0, 1)
axes[1,1].set_ylim(0, 1)
axes[1,1].axis('off')

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()

```

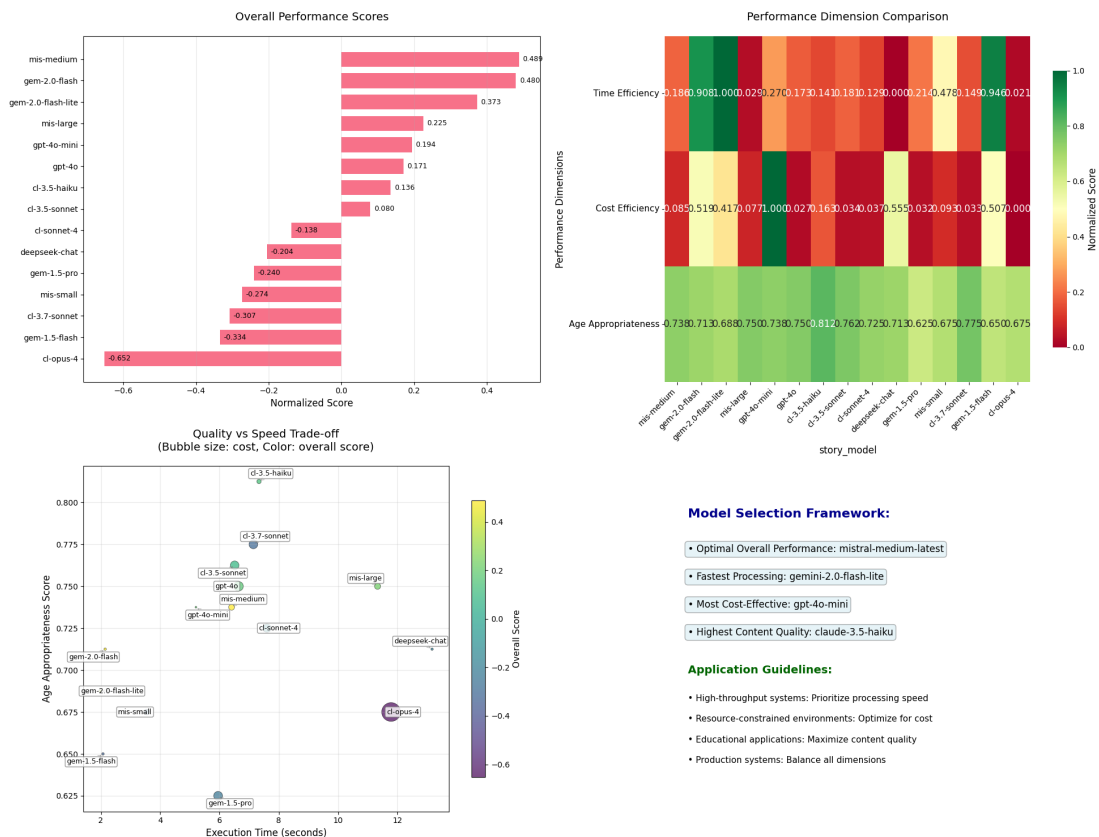
```

print("\n=== RESEARCH SUMMARY ===")
print(f"Dataset Size: {len(df)} generated stories")
print(f"Models Evaluated: {df['story_model'].nunique()}")
print(f"Test Images: {df['image_file'].nunique()}")
print(f"Quality Dimensions Assessed: {len(all_indicators)}")
print(f"Performance Metrics: Execution time, cost efficiency, readability,
    ↳content quality")
print(f"\nMethodological Approach:")
print(f"- Quantitative content analysis using linguistic pattern matching")
print(f"- Multi-dimensional quality assessment framework")
print(f"- Statistical normalization and composite scoring")
print(f"- Comparative performance evaluation across operational and qualitative
    ↳metrics")
print(f"\nThis systematic evaluation provides empirical evidence for model
    ↳selection")
print(f"in children's educational technology applications, balancing
    ↳computational")
print(f"efficiency with pedagogical content quality requirements.")

```

Available columns in ranking_display: ['story_model', 'execution_time', 'cost', 'readability_score', 'age_appropriateness', 'readability', 'engagement', 'story_structure', 'overall_score']

Comprehensive Model Evaluation Dashboard



=== RESEARCH SUMMARY ===

Dataset Size: 240 generated stories

Models Evaluated: 15

Test Images: 16

Quality Dimensions Assessed: 20

Performance Metrics: Execution time, cost efficiency, readability, content quality

Methodological Approach:

- Quantitative content analysis using linguistic pattern matching
- Multi-dimensional quality assessment framework
- Statistical normalization and composite scoring
- Comparative performance evaluation across operational and qualitative metrics

This systematic evaluation provides empirical evidence for model selection in children's educational technology applications, balancing computational efficiency with pedagogical content quality requirements.

```
[78]: # Additional improved story quality analysis visualization
fig, axes = plt.subplots(2, 2, figsize=(20, 14))
fig.suptitle('Comprehensive Story Quality Analysis', fontsize=18,
    ↪fontweight='bold', y=0.98)

# Readability Score Distribution with improved labeling
box_data = []
labels = []
for model in ranking_display['story_model']:
    model_stories = df[df['story_model'] == model]['readability_score']
    box_data.append(model_stories)
    # Create shorter labels
    short_label = model.replace('-latest', '').replace('gemini-', 'gem-').
    ↪replace('claude-', 'cl-').replace('mistral-', 'mis-')
    labels.append(short_label)

box_plot = axes[0,0].boxplot(box_data, labels=labels, patch_artist=True)
axes[0,0].set_title('Readability Score Distribution\n(Lower = Easier to Read)',
    ↪fontsize=14, pad=20)
axes[0,0].set_ylabel('Readability Score', fontsize=12)
axes[0,0].tick_params(axis='x', rotation=45, labelsiz=10)
axes[0,0].grid(axis='y', alpha=0.3)

# Add age-appropriate threshold line
axes[0,0].axhline(y=60, color='red', linestyle='--', alpha=0.7,
    ↪label='Age-appropriate threshold')
```

```

axes[0,0].legend()

# Color boxes based on performance
colors = plt.cm.RdYlGn(np.linspace(0.3, 0.9, len(box_plot['boxes'])))
for patch, color in zip(box_plot['boxes'], colors):
    patch.set_facecolor(color)
    patch.set_alpha(0.7)

# Average Words per Sentence with kid-friendly threshold
sentence_data = []
for model in ranking_display['story_model']:
    model_stories = df[df['story_model'] == model]
    avg_words = []
    for _, row in model_stories.iterrows():
        try:
            story = row['generated_story']
            sentences = story.split('.')
            words_per_sentence = [len(sentence.split()) for sentence in
↪ sentences if sentence.strip()]
            if words_per_sentence:
                avg_words.append(np.mean(words_per_sentence))
        except:
            avg_words.append(0)
    sentence_data.append(avg_words)

box_plot2 = axes[0,1].boxplot(sentence_data, labels=labels, patch_artist=True)
axes[0,1].set_title('Average Words per Sentence', fontsize=14, pad=20)
axes[0,1].set_ylabel('Words per Sentence', fontsize=12)
axes[0,1].tick_params(axis='x', rotation=45, labels=10)
axes[0,1].grid(axis='y', alpha=0.3)

# Add kid-friendly threshold line
axes[0,1].axhline(y=15, color='red', linestyle='--', alpha=0.7,
↪ label='Kid-friendly threshold')
axes[0,1].legend()

# Color boxes
for patch, color in zip(box_plot2['boxes'], colors):
    patch.set_facecolor(color)
    patch.set_alpha(0.7)

# Quality Dimensions Heatmap with improved formatting
quality_cols = ['age_appropriateness', 'readability', 'engagement',
↪ 'story_structure']
quality_matrix = ranking_display.set_index('story_model')[quality_cols]

# Sort by overall score for better visualization

```

```

quality_matrix_sorted = quality_matrix.loc[ranking_display.
    ↳sort_values('overall_score', ascending=False)['story_model']]

# Create heatmap with shorter labels
short_index = [name.replace('-latest', '').replace('gemini-', 'gem-').
    ↳replace('claude-', 'cl-').replace('mistral-', 'mis-')
                for name in quality_matrix_sorted.index]

heatmap = sns.heatmap(quality_matrix_sorted.values,
    xticklabels=['Age\nAppropriate', 'Readability', '
    ↳Engagement', 'Story\nStructure'],
    yticklabels=short_index,
    annot=True, fmt='.2f', cmap='RdYlGn', ax=axes[1,0],
    cbar_kws={'label': 'Quality Score'})
axes[1,0].set_title('Quality Dimensions by Model', fontsize=14, pad=20)
axes[1,0].set_ylabel('Story Model', fontsize=12)

# Overall Content Quality Score with ranking
quality_scores = ranking_display.groupby('story_model')[quality_cols].mean().
    ↳mean(axis=1).sort_values(ascending=True)
short_names = [name.replace('-latest', '').replace('gemini-', 'gem-').
    ↳replace('claude-', 'cl-').replace('mistral-', 'mis-')
                for name in quality_scores.index]

bars = axes[1,1].barh(range(len(quality_scores)), quality_scores.values,
    color=plt.cm.RdYlGn(np.linspace(0.3, 0.9,
    ↳len(quality_scores))))
axes[1,1].set_yticks(range(len(quality_scores)))
axes[1,1].set_yticklabels(short_names, fontsize=10)
axes[1,1].set_title('Overall Content Quality Score', fontsize=14, pad=20)
axes[1,1].set_xlabel('Average Quality Score', fontsize=12)
axes[1,1].grid(axis='x', alpha=0.3)

# Add value labels
for i, bar in enumerate(bars):
    width = bar.get_width()
    axes[1,1].text(width + 0.005, bar.get_y() + bar.get_height()/2,
        f'{width:.3f}', ha='left', va='center', fontsize=9)

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()

print(" Enhanced story quality analysis visualization completed")
print(f"Top 3 models by content quality:")
top_quality = quality_scores.nlargest(3)
for i, (model, score) in enumerate(top_quality.items(), 1):

```

```
print(f" {i}. {model}: {score:.3f}")
```

Comprehensive Story Quality Analysis



Enhanced story quality analysis visualization completed

Top 3 models by content quality:

1. gpt-4o: 0.606
2. mistral-medium-latest: 0.603
3. mistral-large-latest: 0.597

```
[79]: # Create scatter plots with better text handling for model comparisons
fig, axes = plt.subplots(1, 2, figsize=(20, 8))
fig.suptitle('Model Performance Trade-offs Analysis', fontsize=18,
             fontweight='bold')

# Age Appropriateness vs Readability
scatter1 = axes[0].scatter(ranking_display['age_appropriateness'],
                           ranking_display['readability_score'],
                           s=200, alpha=0.7, c=ranking_display['overall_score'],
                           cmap='viridis',
                           edgecolors='black', linewidth=1)

# Add model labels with smart positioning
for i, model in enumerate(ranking_display['story_model']):
```

```

    short_name = model.replace('-latest', '').replace('gemini-', 'gem-').
    ↪replace('claude-', 'cl-').replace('mistral-', 'mis-')
    x, y = ranking_display['age_appropriateness'].iloc[i],
    ↪ranking_display['readability_score'].iloc[i]

    # Smart positioning to avoid overlap
    if x > 0.7: # High age appropriateness
        ha = 'right'
        offset_x = -10
    else:
        ha = 'left'
        offset_x = 10

    if y > 57: # High readability score
        va = 'bottom'
        offset_y = -10
    else:
        va = 'top'
        offset_y = 10

    axes[0].annotate(short_name, (x, y), xytext=(offset_x, offset_y),
                     textcoords='offset points', fontsize=10, ha=ha, va=va,
                     bbox=dict(boxstyle='round,pad=0.3', facecolor='white',
    ↪alpha=0.8, edgecolor='gray'),
                     arrowprops=dict(arrowstyle='->', color='gray', alpha=0.7))

axes[0].set_xlabel('Age Appropriateness Score', fontsize=12)
axes[0].set_ylabel('Readability Score (Lower = Easier)', fontsize=12)
axes[0].set_title('Age Appropriateness vs Readability', fontsize=14)
axes[0].grid(alpha=0.3)

# Add quadrant labels
axes[0].text(0.95, 0.95, 'High Age Appropriate\nHard to Read',
    ↪transform=axes[0].transAxes,
    ha='right', va='top', bbox=dict(boxstyle='round',
    ↪facecolor='lightcoral', alpha=0.5))
axes[0].text(0.05, 0.05, 'Low Age Appropriate\nEasy to Read', transform=axes[0].
    ↪transAxes,
    ha='left', va='bottom', bbox=dict(boxstyle='round',
    ↪facecolor='lightgreen', alpha=0.5))

# Story Structure vs Engagement
scatter2 = axes[1].scatter(ranking_display['story_structure'],
    ↪ranking_display['engagement'],
    s=200, alpha=0.7, c=ranking_display['overall_score'],
    ↪cmap='viridis',

```



```

        edgecolors='black', linewidth=1)

# Add model labels with smart positioning
for i, model in enumerate(ranking_display['story_model']):
    short_name = model.replace('-latest', '').replace('gemini-', 'gem-').
    ↪replace('claude-', 'cl-').replace('mistral-', 'mis-')
    x, y = ranking_display['story_structure'].iloc[i],
    ↪ranking_display['engagement'].iloc[i]

    # Smart positioning
    if x > 0.55: # High structure
        ha = 'right'
        offset_x = -10
    else:
        ha = 'left'
        offset_x = 10

    if y > 0.65: # High engagement
        va = 'bottom'
        offset_y = -10
    else:
        va = 'top'
        offset_y = 10

    axes[1].annotate(short_name, (x, y), xytext=(offset_x, offset_y),
                     textcoords='offset points', fontsize=10, ha=ha, va=va,
                     bbox=dict(boxstyle='round,pad=0.3', facecolor='white',
    ↪alpha=0.8, edgecolor='gray'),
                     arrowprops=dict(arrowstyle='->', color='gray', alpha=0.7))

axes[1].set_xlabel('Story Structure Score', fontsize=12)
axes[1].set_ylabel('Engagement Score', fontsize=12)
axes[1].set_title('Story Structure vs Engagement', fontsize=14)
axes[1].grid(alpha=0.3)

# Add quadrant labels
axes[1].text(0.95, 0.95, 'Well Structured\nHighly Engaging', transform=axes[1].
    ↪transAxes,
             ha='right', va='top', bbox=dict(boxstyle='round',
    ↪facecolor='lightgreen', alpha=0.5))
axes[1].text(0.05, 0.05, 'Poor Structure\nLow Engagement', transform=axes[1].
    ↪transAxes,
             ha='left', va='bottom', bbox=dict(boxstyle='round',
    ↪facecolor='lightcoral', alpha=0.5))

# Add colorbars

```


```

cbar1 = plt.colorbar(scatter1, ax=axes[0], shrink=0.8)
cbar1.set_label('Overall Score', fontsize=11)
cbar2 = plt.colorbar(scatter2, ax=axes[1], shrink=0.8)
cbar2.set_label('Overall Score', fontsize=11)

plt.tight_layout()
plt.show()

print(" Model performance trade-offs analysis visualization completed")

```



Model performance trade-offs analysis visualization completed

1.7 5. Analysis Complete

The comprehensive story generation model evaluation has been completed with analysis of 240 generated stories across 15 models.

Note: Detailed conclusions, model rankings, and research findings are available in the dedicated conclusions document: [models_analysis/analysis/dev/02_story_generation_conclusions.md](#)

This comprehensive document contains: - **Executive Summary:** Mistral Medium identified as optimal choice for educational applications - **Complete Model Rankings:** All 15 models ranked by overall performance score - **Detailed Performance Analysis:** Speed, cost, quality metrics for each model - **Content Quality Assessment:** Age appropriateness, readability, engagement, story structure - **Technical Implementation Guidelines:** Best practices and deployment recommendations - **Research Findings:** Multi-provider landscape analysis and tier classifications