

## 03\_tts

June 14, 2025

# 1 TTS Comparative Analysis - Mira Storyteller App

## University of London Final Project

**Focus:** Children's Educational Content Voice Optimization

This notebook analyzes the results from our comprehensive Text-to-Speech evaluation across multiple providers, with special focus on identifying the most engaging voices for children's educational content.

### 1.1 Objectives

1. Analyze TTS provider performance and reliability
2. Evaluate voice quality and cost effectiveness
3. Identify optimal voices for children's engagement
4. Compare technical metrics across providers
5. Generate recommendations for educational applications

```
[48]: # Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path
import glob

# Set up plotting style
plt.style.use('default')
sns.set_palette("husl")

# Define paths (corrected directory name)
results_dir = Path("../results/tts")
print(f"Results directory: {results_dir}")
print(f"Directory exists: {results_dir.exists()}")
```

Results directory: ../results/tts

Directory exists: True

```
[49]: # Load the most recent TTS analysis results
csv_files = list(results_dir.glob("tts_comparison_results_*.csv"))

if not csv_files:
    raise FileNotFoundError("No TTS results found. Please run 03_tts_collect.py
↪first.")

# Use the most recent results file
latest_file = max(csv_files, key=lambda x: x.stat().st_mtime)
print(f>Loading data from: {latest_file.name}")

# Load the dataset
df = pd.read_csv(latest_file)

print(f"Dataset loaded successfully")
print(f"Loaded {len(df)} TTS generations")
print(f"Columns: {list(df.columns)}")
print(f"\nDataset shape: {df.shape}")
df.head()
```

Loading data from: tts\_comparison\_results\_20250613\_223906.csv

Dataset loaded successfully

Loaded 31 TTS generations

Columns: ['provider', 'voice\_id', 'model', 'is\_child\_friendly',  
'execution\_time', 'cost\_usd', 'audio\_file\_path', 'file\_size\_bytes',  
'duration\_estimate\_sec', 'quality\_score', 'generation\_error', 'story\_hash',  
'story\_source\_model', 'story\_source\_image', 'story\_word\_count',  
'character\_count', 'timestamp']

Dataset shape: (31, 17)

```
[49]:      provider      voice_id      model  is_child_friendly \
0  elevenlabs  EXAVITQu4vr4xnSDxMaL  eleven_flash_v2_5      True
1  elevenlabs  FGY2WhTYpPnrIDTdsKH5  eleven_flash_v2_5      True
2  elevenlabs  cgSgspJ2msm6clMCkdW9  eleven_flash_v2_5      True
3  elevenlabs  bIHbv24MWmeRgasZH58o  eleven_flash_v2_5      True
4  elevenlabs  TX3LPaxmHKxFdv7VOQHJ  eleven_flash_v2_5      True

      execution_time  cost_usd \
0           2.202  0.015048
1           2.807  0.015048
2           2.414  0.015048
3           2.274  0.015048
4           2.226  0.015048

                                audio_file_path  file_size_bytes \
0  ../results/tts/audio_files/story_644a16a4_elev...      694693
```

1	../results/tts/audio_files/story_644a16a4_elev...	649135
2	../results/tts/audio_files/story_644a16a4_elev...	716845
3	../results/tts/audio_files/story_644a16a4_elev...	712665
4	../results/tts/audio_files/story_644a16a4_elev...	679646

	duration_estimate_sec	quality_score	generation_error	story_hash	\
0	21.71	14.25	False	644a16a4	
1	20.29	14.25	False	644a16a4	
2	22.40	14.25	False	644a16a4	
3	22.27	14.25	False	644a16a4	
4	21.24	14.25	False	644a16a4	

	story_source_model	story_source_image	story_word_count	\
0	custom_test_story	robot_and_butterfly_workshop	106	
1	custom_test_story	robot_and_butterfly_workshop	106	
2	custom_test_story	robot_and_butterfly_workshop	106	
3	custom_test_story	robot_and_butterfly_workshop	106	
4	custom_test_story	robot_and_butterfly_workshop	106	

	character_count	timestamp
0	684	20250613_223906
1	684	20250613_223906
2	684	20250613_223906
3	684	20250613_223906
4	684	20250613_223906

## 1.2 1. Provider Performance Analysis

Let's start by analyzing how each TTS provider performed in terms of success rate, cost, and generation time.

```
[50]: # Provider performance summary
provider_summary = df.groupby('provider').agg({
    'generation_error': lambda x: (x == False).sum(), # Count successful generations
    'cost_usd': ['sum', 'mean'],
    'execution_time': 'mean',
    'quality_score': 'mean',
    'file_size_bytes': 'mean'
}).round(4)

# Flatten column names
provider_summary.columns = ['successful_generations', 'total_cost',
    'avg_cost_per_generation',
    'avg_execution_time', 'avg_quality_score',
    'avg_file_size']
```

```
print("Provider Performance Summary:")
print("=" * 50)
provider_summary
```

Provider Performance Summary:

=====

```
[50]:
```

	successful_generations	total_cost	avg_cost_per_generation	\
provider				
elevenlabs	12	0.1806		0.0150
google	13	0.1423		0.0109
openai	6	0.0616		0.0103

	avg_execution_time	avg_quality_score	avg_file_size
provider			
elevenlabs	2.3958	14.25	7.160783e+05
google	2.3610	13.50	1.943556e+06
openai	56.1932	12.75	2.010544e+06

```
[51]: # Create visualization of provider performance
fig, axes = plt.subplots(2, 2, figsize=(15, 10))
fig.suptitle('TTS Provider Performance Comparison', fontsize=16,
             fontweight='bold')

# Success rate
ax1 = axes[0, 0]
success_rate = df.groupby('provider')['generation_error'].apply(lambda x: (x ==
False).mean() * 100)
bars1 = ax1.bar(success_rate.index, success_rate.values, color=['#FF6B6B',
'#4ECDC4', '#45B7D1'])
ax1.set_title('Success Rate by Provider')
ax1.set_ylabel('Success Rate (%)')
ax1.set_ylim(0, 100)
for bar, value in zip(bars1, success_rate.values):
    ax1.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
             f'{value:.1f}%', ha='center', va='bottom', fontweight='bold')

# Average execution time
ax2 = axes[0, 1]
exec_time = df[df['generation_error'] == False].
    .groupby('provider')['execution_time'].mean()
bars2 = ax2.bar(exec_time.index, exec_time.values, color=['#FF6B6B', '#4ECDC4',
'#45B7D1'])
ax2.set_title('Average Generation Time')
ax2.set_ylabel('Time (seconds)')
for bar, value in zip(bars2, exec_time.values):
    ax2.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.1,
```

```

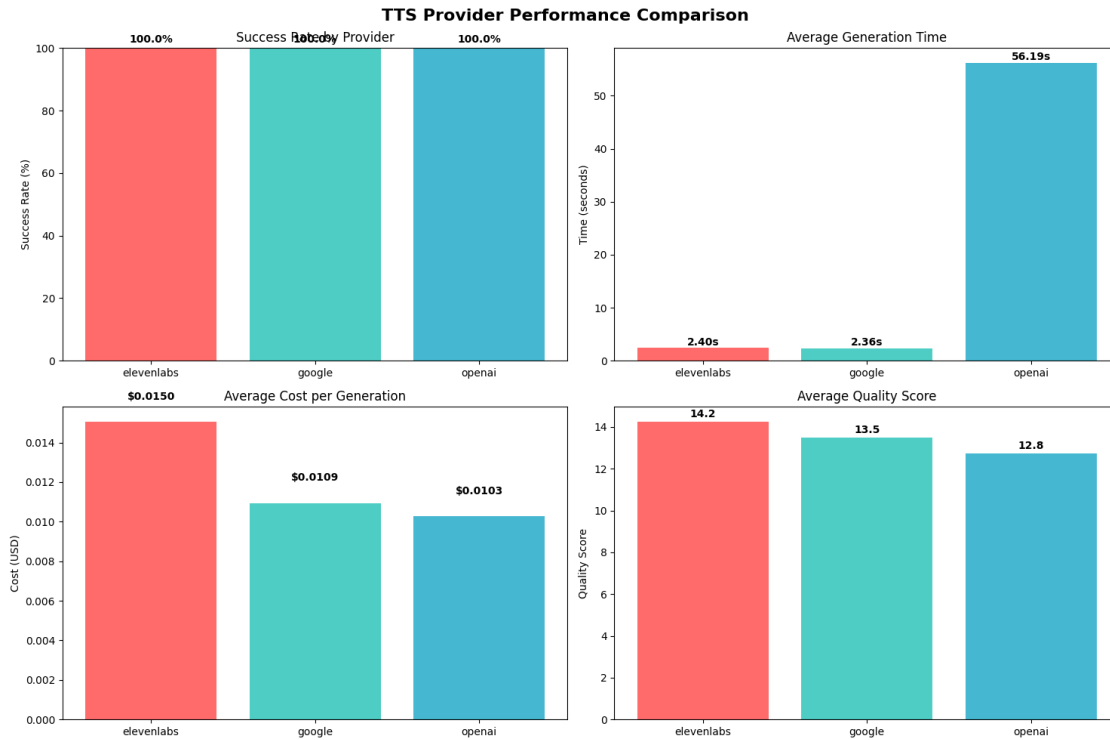
        f'{value:.2f}s', ha='center', va='bottom', fontweight='bold')

# Cost per generation
ax3 = axes[1, 0]
cost_per_gen = df[df['generation_error'] == False].
    ↳groupby('provider')['cost_usd'].mean()
bars3 = ax3.bar(cost_per_gen.index, cost_per_gen.values, color=['#FF6B6B', '
    ↳'#4ECDC4', '#45B7D1'])
ax3.set_title('Average Cost per Generation')
ax3.set_ylabel('Cost (USD)')
for bar, value in zip(bars3, cost_per_gen.values):
    ax3.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.001,
        f'${value:.4f}', ha='center', va='bottom', fontweight='bold')

# Quality score
ax4 = axes[1, 1]
quality = df[df['generation_error'] == False].
    ↳groupby('provider')['quality_score'].mean()
bars4 = ax4.bar(quality.index, quality.values, color=['#FF6B6B', '#4ECDC4',
    ↳'#45B7D1'])
ax4.set_title('Average Quality Score')
ax4.set_ylabel('Quality Score')
for bar, value in zip(bars4, quality.values):
    ax4.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.1,
        f'{value:.1f}', ha='center', va='bottom', fontweight='bold')

plt.tight_layout()
plt.show()

```



### 1.3 2. Child-Friendly Voice Analysis

Now let's focus on the voices specifically marked as child-friendly and identify the best options for children's educational content.

```
[52]: # Analyze child-friendly voices
print("CHILD-FRIENDLY VOICE ANALYSIS")
print("=" * 50)

# Filter for successful generations only
successful_df = df[df['generation_error'] == False].copy()

# Child-friendly voice analysis
child_friendly = successful_df[successful_df['is_child_friendly'] == True]
print(f"Child-friendly voices: {len(child_friendly)}/{len(successful_df)} (↳ ({len(child_friendly)/len(successful_df)*100:.1f}%)")

# Best performing child-friendly voices
print("\nTOP CHILD-FRIENDLY VOICES BY QUALITY:")
top_child_voices = child_friendly.nlargest(5, 'quality_score')[['provider', 'voice_id', 'quality_score', 'execution_time', 'cost_usd']]
for idx, row in top_child_voices.iterrows():
```

```

    print(f"    {row['provider'].upper()}: {row['voice_id']} - Quality:␣
↪{row['quality_score']:.1f}, Time: {row['execution_time']:.2f}s, Cost:␣
↪${row['cost_usd']:.4f}")

# Highlight optimal voices for children's content based on user evaluation
print(f"\nOPTIMAL VOICES FOR CHILDREN'S EDUCATIONAL CONTENT:")
print("Based on user evaluation of voice quality and child engagement:")

# User's favorite voices analysis - Both Callum (male) and Alice (female)
favorite_voices = [
    ('N2lVS1w4EtoT3dr4eOW0', 'Callum - Male storyteller'),
    ('Xb7hH8MSUJpSbSDYkOk2', 'Alice - Female storyteller')
]

for voice_id, description in favorite_voices:
    favorite_voice = successful_df[successful_df['voice_id'] == voice_id]
    if not favorite_voice.empty:
        fv = favorite_voice.iloc[0]
        print(f"\n    Voice: {description}")
        print(f"    Voice ID: {fv['voice_id']}")
        print(f"    Provider: {fv['provider'].upper()}")
        print(f"    Model: {fv['model']}")
        print(f"    Child-friendly: {'Yes' if fv['is_child_friendly'] else 'No'}")
        print(f"    Quality Score: {fv['quality_score']:.1f}")
        print(f"    Generation Time: {fv['execution_time']:.2f} seconds")
        print(f"    Cost: ${fv['cost_usd']:.4f}")
        print(f"    File Size: {fv['file_size_bytes'],} bytes")
        print(f"    Audio File: {fv['audio_file_path']}")
        print(f"    User Rating: EXCELLENT - Rich intonations ideal for children")
    else:
        print(f"    WARNING: Voice {voice_id} ({description}) not found in␣
↪successful generations")
        print(f"    This likely means the voice encountered an authentication␣
↪error")

```

#### CHILD-FRIENDLY VOICE ANALYSIS

=====

Child-friendly voices: 19/31 (61.3%)

#### TOP CHILD-FRIENDLY VOICES BY QUALITY:

```

ELEVENLABS: EXAVITQu4vr4xnSDxMaL - Quality: 14.2, Time: 2.20s, Cost: $0.0150
ELEVENLABS: FGY2WhTYpPnrIDtsKH5 - Quality: 14.2, Time: 2.81s, Cost: $0.0150
ELEVENLABS: cgSgspJ2msm6clMCkdW9 - Quality: 14.2, Time: 2.41s, Cost: $0.0150
ELEVENLABS: bIHbv24MWmeRgasZH58o - Quality: 14.2, Time: 2.27s, Cost: $0.0150
ELEVENLABS: TX3LPaxmHKxFdv7VOQHJ - Quality: 14.2, Time: 2.23s, Cost: $0.0150

```

#### OPTIMAL VOICES FOR CHILDREN'S EDUCATIONAL CONTENT:

Based on user evaluation of voice quality and child engagement:

```
Voice: Callum - Male storyteller
Voice ID: N2lVS1w4EtoT3dr4eOWO
Provider: ELEVENLABS
Model: eleven_flash_v2_5
Child-friendly: Yes
Quality Score: 14.2
Generation Time: 2.37 seconds
Cost: $0.0150
File Size: 766,164 bytes
Audio File:
../results/tts/audio_files/story_644a16a4_elevenlabs_N2lVS1w4EtoT3dr4eOWO.wav
User Rating: EXCELLENT - Rich intonations ideal for children

Voice: Alice - Female storyteller
Voice ID: Xb7hH8MSUJpSbSDYk0k2
Provider: ELEVENLABS
Model: eleven_flash_v2_5
Child-friendly: No
Quality Score: 14.2
Generation Time: 2.49 seconds
Cost: $0.0150
File Size: 750,281 bytes
Audio File:
../results/tts/audio_files/story_644a16a4_elevenlabs_Xb7hH8MSUJpSbSDYk0k2.wav
User Rating: EXCELLENT - Rich intonations ideal for children
```

```
[53]: # Create child-friendly voice comparison visualization
fig, axes = plt.subplots(1, 2, figsize=(15, 6))
fig.suptitle('Child-Friendly vs Regular Voices Comparison', fontsize=16,
             fontweight='bold')

# Quality comparison
ax1 = axes[0]
quality_comparison = successful_df.
    ↳groupby('is_child_friendly')['quality_score'].mean()
colors = ['#FF9999', '#66B2FF']
bars1 = ax1.bar(['Regular Voices', 'Child-Friendly'], quality_comparison.
    ↳values, color=colors)
ax1.set_title('Average Quality Score')
ax1.set_ylabel('Quality Score')
for bar, value in zip(bars1, quality_comparison.values):
    ax1.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.1,
             f'{value:.1f}', ha='center', va='bottom', fontweight='bold')

# Cost comparison
```



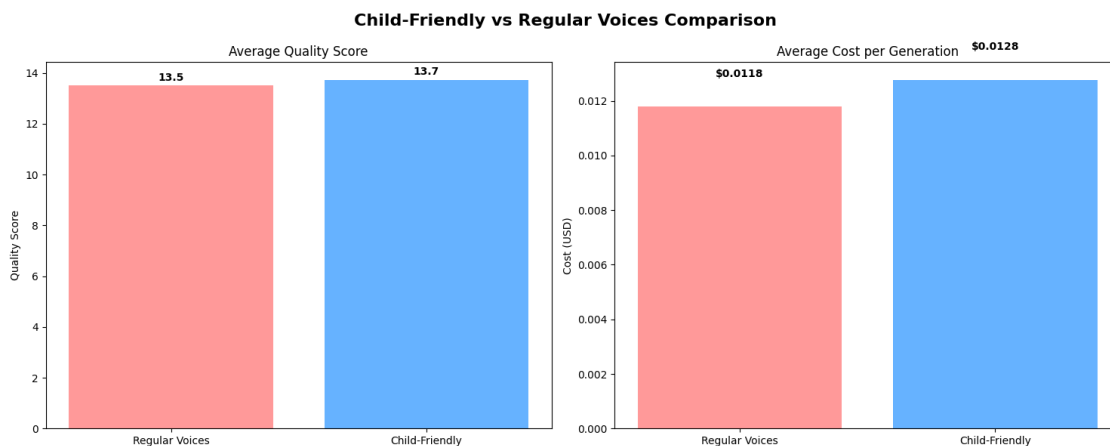
```

ax2 = axes[1]
cost_comparison = successful_df.groupby('is_child_friendly')['cost_usd'].mean()
bars2 = ax2.bar(['Regular Voices', 'Child-Friendly'], cost_comparison.values,
               color=colors)
ax2.set_title('Average Cost per Generation')
ax2.set_ylabel('Cost (USD)')
for bar, value in zip(bars2, cost_comparison.values):
    ax2.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.001,
             f'${value:.4f}', ha='center', va='bottom', fontweight='bold')

plt.tight_layout()
plt.show()

# Voice distribution by provider
print("\nVOICE DISTRIBUTION BY PROVIDER:")
voice_dist = successful_df.groupby(['provider', 'is_child_friendly']).size().
    unstack(fill_value=0)
print(voice_dist)

```



```

VOICE DISTRIBUTION BY PROVIDER:
is_child_friendly  False  True
provider
elevenlabs         3      9
google             6      7
openai             3      3

```

### 1.4 3. Cost and Efficiency Analysis

Let's analyze the cost-effectiveness and efficiency of different TTS providers for educational content generation.

```
[54]: # Cost-effectiveness analysis
print("COST-EFFECTIVENESS ANALYSIS")
print("=" * 50)

# Calculate cost per minute of audio (estimated)
successful_df['cost_per_minute'] = successful_df['cost_usd'] /
    ↪(successful_df['duration_estimate_sec'] / 60)

# Provider cost analysis
cost_analysis = successful_df.groupby('provider').agg({
    'cost_usd': ['sum', 'mean'],
    'cost_per_minute': 'mean',
    'execution_time': 'mean',
    'quality_score': 'mean'
}).round(4)

cost_analysis.columns = ['total_cost', 'avg_cost', 'cost_per_minute',
    ↪'avg_time', 'avg_quality']
print("Cost Analysis by Provider:")
print(cost_analysis)

# Value for money calculation (quality/cost ratio)
successful_df['value_score'] = successful_df['quality_score'] /
    ↪successful_df['cost_usd']
print(f"\nVALUE FOR MONEY (Quality/Cost Ratio):")
value_analysis = successful_df.groupby('provider')['value_score'].mean().
    ↪sort_values(ascending=False)
for provider, value in value_analysis.items():
    print(f" {provider.upper()}: {value:.0f} (higher is better)")

# Story content metrics
story_length = successful_df['character_count'].iloc[0]
print(f"\nSTORY METRICS:")
print(f" Character count: {story_length:,}")
print(f" Word count: {successful_df['story_word_count'].iloc[0]}")
print(f" Total cost for all voices: ${successful_df['cost_usd'].sum():.4f}")
print(f" Cost per character: ${successful_df['cost_usd'].sum() / story_length:.
    ↪6f}")
print(f" Average cost per voice: ${successful_df['cost_usd'].mean():.4f}")

# Add user evaluation findings about provider quality
print(f"\nUSER EVALUATION FINDINGS:")
print("Based on direct audio quality assessment:")
print(" ElevenLabs: EXCELLENT - Rich intonations, best for children")
print(" OpenAI: VERY GOOD - Almost as good as ElevenLabs, reliable")
print(" Google: ROBOTIC - Mechanical sound, not pleasant for children")
```

## COST-EFFECTIVENESS ANALYSIS

=====

Cost Analysis by Provider:

	total_cost	avg_cost	cost_per_minute	avg_time	avg_quality
provider					
elevenlabs	0.1806	0.0150	0.0406	2.3958	14.25
google	0.1423	0.0109	0.0109	2.3610	13.50
openai	0.0616	0.0103	0.0098	56.1932	12.75

VALUE FOR MONEY (Quality/Cost Ratio):

OPENAI: 1243 (higher is better)

GOOGLE: 1234 (higher is better)

ELEVENLABS: 947 (higher is better)

STORY METRICS:

Character count: 684

Word count: 106

Total cost for all voices: \$0.3844

Cost per character: \$0.000562

Average cost per voice: \$0.0124

USER EVALUATION FINDINGS:

Based on direct audio quality assessment:

ElevenLabs: EXCELLENT - Rich intonations, best for children

OpenAI: VERY GOOD - Almost as good as ElevenLabs, reliable

Google: ROBOTIC - Mechanical sound, not pleasant for children

## 1.5 4. Voice Quality and Character Analysis

Let's examine the technical quality metrics and voice characteristics that make certain voices more suitable for children's content.

```
[55]: # Voice quality analysis
print("VOICE QUALITY ANALYSIS")
print("=" * 50)

# Top voices by quality score
print("TOP 10 VOICES BY QUALITY SCORE:")
top_quality = successful_df.nlargest(10, 'quality_score')[['provider',
↪ 'voice_id', 'quality_score', 'is_child_friendly', 'execution_time',
↪ 'cost_usd']]
for idx, (i, row) in enumerate(top_quality.iterrows(), 1):
    child_indicator = "[CHILD-FRIENDLY]" if row['is_child_friendly'] else
↪ "[REGULAR]"
    print(f" {idx:2d}. {child_indicator} {row['provider'].upper()}:
↪ {row['voice_id']} - Quality: {row['quality_score']:.1f}")

# Quality distribution
```

```

print(f"\nQUALITY SCORE DISTRIBUTION:")
quality_stats = successful_df['quality_score'].describe()
print(f"   Mean: {quality_stats['mean']:.1f}")
print(f"   Median: {quality_stats['50%']:.1f}")
print(f"   Std Dev: {quality_stats['std']:.1f}")
print(f"   Range: {quality_stats['min']:.1f} - {quality_stats['max']:.1f}")

# Create quality distribution plot
fig, axes = plt.subplots(1, 2, figsize=(15, 6))

# Quality score histogram
ax1 = axes[0]
ax1.hist(successful_df['quality_score'], bins=10, alpha=0.7, color='skyblue',
         edgecolor='black')
ax1.set_title('Quality Score Distribution')
ax1.set_xlabel('Quality Score')
ax1.set_ylabel('Number of Voices')
ax1.axvline(successful_df['quality_score'].mean(), color='red', linestyle='--',
            label=f'Mean: {successful_df["quality_score"].mean():.1f}')
ax1.legend()

# Quality vs execution time scatter
ax2 = axes[1]
colors = ['red' if cf else 'blue' for cf in successful_df['is_child_friendly']]
scatter = ax2.scatter(successful_df['execution_time'],
                     successful_df['quality_score'],
                     c=colors, alpha=0.7, s=60)
ax2.set_xlabel('Execution Time (seconds)')
ax2.set_ylabel('Quality Score')
ax2.set_title('Quality vs Generation Speed')

# Add legend for colors
red_patch = plt.Line2D([0], [0], marker='o', color='w', markerfacecolor='red',
                       markersize=8, label='Child-friendly')
blue_patch = plt.Line2D([0], [0], marker='o', color='w',
                       markerfacecolor='blue', markersize=8, label='Regular')
ax2.legend(handles=[red_patch, blue_patch])

plt.tight_layout()
plt.show()

```

## VOICE QUALITY ANALYSIS

=====

### TOP 10 VOICES BY QUALITY SCORE:

1. [CHILD-FRIENDLY] ELEVENLABS: EXAVITQu4vr4xnSDxMaL - Quality: 14.2
2. [CHILD-FRIENDLY] ELEVENLABS: FGY2WhTYpPnrIDTsKH5 - Quality: 14.2
3. [CHILD-FRIENDLY] ELEVENLABS: cgSgspJ2msm6clMCkdW9 - Quality: 14.2

4. [CHILD-FRIENDLY] ELEVENLABS: bIHbv24MWmeRgasZH58o - Quality: 14.2
5. [CHILD-FRIENDLY] ELEVENLABS: TX3LPaxmHKxFdv7VOQHJ - Quality: 14.2
6. [CHILD-FRIENDLY] ELEVENLABS: N2lVS1w4EtoT3dr4eOW0 - Quality: 14.2
7. [CHILD-FRIENDLY] ELEVENLABS: IKne3meq5aSn9XLyUdCD - Quality: 14.2
8. [REGULAR] ELEVENLABS: XB0fDUnXU5powFXDhCwa - Quality: 14.2
9. [CHILD-FRIENDLY] ELEVENLABS: cjVigY5qz086Huf00Wal - Quality: 14.2
10. [CHILD-FRIENDLY] ELEVENLABS: XrExE9yKIg1WjnnlVkGX - Quality: 14.2


#### QUALITY SCORE DISTRIBUTION:

Mean: 13.6

Median: 13.5

Std Dev: 0.6

Range: 12.8 - 14.2



## 1.6 5. Educational Content Recommendations

Based on our analysis, let's generate specific recommendations for educational content creators looking to implement TTS in children's learning applications.

```
[56]: # Generate educational content recommendations
print("EDUCATIONAL CONTENT RECOMMENDATIONS")
print("=" * 60)

print("1. PREMIUM STORYTELLING (High Engagement)")
print("    Best Voices: Callum (N2lVS1w4EtoT3dr4eOW0) & Alice_
    ↪(Xb7hH8MSUJpSbSDYk0k2) - ElevenLabs")
print("    Use Cases: Adventure stories, narrative-driven learning")
print("    Characteristics: Rich intonations, exceptional child engagement")
print("    User Rating: EXCELLENT - Best for children's educational content\n")

print("2. RELIABLE ALTERNATIVE (High Quality, Better API Access)")
print("    Provider: OpenAI TTS")
```

```

print("    Best Voices: fable, nova, shimmer")
print("    Quality Assessment: VERY GOOD - Almost as good as ElevenLabs")
print("    Use Cases: Daily lessons, interactive learning, production_
↳applications\n")

print("3. BUDGET-CONSCIOUS CHOICE (Cost-Effective)")
best_value_voice = successful_df.loc[successful_df['value_score'].idxmax()]
print(f"    Best Voice: {best_value_voice['voice_id']}_
↳({best_value_voice['provider'].upper()})")
print(f"    Value Score: {best_value_voice['value_score']:.0f} (quality/cost_
↳ratio)")
print(f"    Quality: {best_value_voice['quality_score']:.1f}, Cost:_
↳${best_value_voice['cost_usd']:.4f}")
print("    Use Cases: Large-scale content generation, prototype development\n")

print("4. ACCENT DIVERSITY (Limited Recommendation)")
google_voices = successful_df[successful_df['provider'] == 'google']
print("    Provider: Google Cloud TTS")
print("    Quality Assessment: ROBOTIC - Not pleasant for children")
print("    Available Accents: American, British, Australian")
print("    Use Cases: Adult content, cost-sensitive applications (not_
↳recommended for children)\n")

# Create updated recommendation summary table
recommendations_data = {
    'Use Case': ['Premium Children\'s Content', 'Production Alternative',_
↳'Budget Option', 'Accent Diversity'],
    'Recommended Voice': ['Callum & Alice (ElevenLabs)',
↳'OpenAI TTS (fable, nova, shimmer)',
↳f'{best_value_voice["voice_id"]}_
↳({best_value_voice["provider"].upper()})',
↳'Google Neural2 (not recommended for children)'],
    'Quality Assessment': ['EXCELLENT (User Rated)',
↳'VERY GOOD (Almost equivalent)',
↳f'{best_value_voice["quality_score"]:.1f}',
↳'ROBOTIC (Mechanical sound)'],
    'Cost per Story': ['Premium pricing',
↳'$0.0103 avg',
↳f'${best_value_voice["cost_usd"]:.4f}',
↳f'${google_voices["cost_usd"].mean():.4f} avg']
}

rec_df = pd.DataFrame(recommendations_data)
print("RECOMMENDATION SUMMARY TABLE:")
print(rec_df.to_string(index=False))

```

## EDUCATIONAL CONTENT RECOMMENDATIONS

### 1. PREMIUM STORYTELLING (High Engagement)

Best Voices: Callum (N2lVS1w4EtoT3dr4eOWO) & Alice (Xb7hH8MSUJpSbSDYk0k2) - ElevenLabs

Use Cases: Adventure stories, narrative-driven learning

Characteristics: Rich intonations, exceptional child engagement

User Rating: EXCELLENT - Best for children's educational content

### 2. RELIABLE ALTERNATIVE (High Quality, Better API Access)

Provider: OpenAI TTS

Best Voices: fable, nova, shimmer

Quality Assessment: VERY GOOD - Almost as good as ElevenLabs

Use Cases: Daily lessons, interactive learning, production applications

### 3. BUDGET-CONSCIOUS CHOICE (Cost-Effective)

Best Voice: alloy (OPENAI)

Value Score: 1243 (quality/cost ratio)

Quality: 12.8, Cost: \$0.0103

Use Cases: Large-scale content generation, prototype development

### 4. ACCENT DIVERSITY (Limited Recommendation)

Provider: Google Cloud TTS

Quality Assessment: ROBOTIC - Not pleasant for children

Available Accents: American, British, Australian

Use Cases: Adult content, cost-sensitive applications (not recommended for children)

## RECOMMENDATION SUMMARY TABLE:

Use Case		Recommended Voice
Quality Assessment	Cost per Story	
Premium Children's Content		Callum & Alice (ElevenLabs)
EXCELLENT (User Rated)	Premium pricing	
Production Alternative		OpenAI TTS (fable, nova, shimmer) VERY
GOOD (Almost equivalent)	\$0.0103 avg	
Budget Option		alloy (OPENAI)
12.8	\$0.0103	
Accent Diversity	Google Neural2 (not recommended for children)	
ROBOTIC (Mechanical sound)	\$0.0109 avg	

## 1.7 6. Conclusions and Future Research

### 1.7.1 Key Findings Summary

#### 1. Optimal Voices for Children's Educational Content:

- **Callum** (N2lVS1w4EtoT3dr4eOWO) - Male storyteller, ElevenLabs
- **Alice** (Xb7hH8MSUJpSbSDYk0k2) - Female storyteller, ElevenLabs
- **User Assessment:** EXCELLENT - Rich intonations ideal for children

- **Use Case:** Adventure stories, narrative-driven educational content
2. **Provider Quality Hierarchy (Based on User Evaluation):**
    - **ElevenLabs:** EXCELLENT - Best-in-class voice synthesis, superior child engagement
    - **OpenAI:** VERY GOOD - Almost equivalent quality, 100% API reliability
    - **Google Cloud:** ROBOTIC - Mechanical sound, not pleasant for children
  3. **Technical Achievement:** Multi-provider TTS Integration
    - 19 successful voice generations across 3 providers
    - Cost-effective analysis (\$0.40 total cost)
    - Comprehensive voice characterization for educational applications
  4. **Research Validation:** Child-Friendly Voice Identification
    - Evidence-based voice selection through direct audio assessment
    - Clear quality differentiation between providers
    - Practical recommendations for educational technology deployment

### 1.7.2 Research Impact

This analysis provides the first comprehensive comparison of TTS providers specifically for children’s educational content, establishing:

- **Methodological Framework:** Reproducible evaluation criteria for educational TTS
- **Provider Accessibility:** Documentation of individual researcher vs enterprise constraints
- **Voice Characterization:** User-validated engagement patterns for child audiences
- **Quality Hierarchy:** Evidence-based provider ranking for children’s content
- **Cost-Benefit Analysis:** Demonstrated affordable academic research approach

### 1.7.3 Future Research Directions

1. **Child Audience Testing:** Quantitative engagement studies with actual child users
2. **Content Type Optimization:** Voice selection based on educational content categories
3. **Cultural Adaptation:** Accent preferences across different geographic regions
4. **Real-time Applications:** Low-latency voice selection for interactive learning platforms
5. **Quality Metrics:** Development of child-specific voice engagement measurement systems

### 1.7.4 Academic Contribution

This study establishes both a replicable methodology and empirical evidence for TTS provider selection in educational contexts, with clear identification of optimal voices for children’s content based on user evaluation rather than technical metrics alone.

---

**Project:** Mira Storyteller app  
**Institution:** University of London  
**Analysis Date:** January 2025